

cerned entirely with statistical evidence about error rates in capital cases, an issue that, of course, bears importantly on the morality of the death penalty. And in a few other cases, the readings do not clearly proceed from some moral theory. So, not every article summary includes a recommendation to consult the moral theory primer. But most of the reading selections do connect directly with one or more of the moral theories explained in the primer.

- *Quick reference guide to moral theories.* In order to make it easy to review the fundamental principles of each of the theories, I have included an appendix, “Quick Guide to Moral Theories” with this information. Once one has read the relevant sections of the moral theory primer, this guide may be consulted to refresh one’s memory of the basics.

Again, the preceding steps reflect how I like to proceed. Users are invited to find ways that best fit their own style of teaching.

NOTE

1. Of course, some topics naturally go well together because the moral issues they raise are deeply connected. For instance, chapter 3 on pornography, hate speech, and censorship raises issues about the morality of government interference in the lives of its citizens. The same sort of issue comes up in chapter 4 on drugs, gambling, and addiction. Chapters 12 and 13 on animals and the environment, respectively, go together because they raise important questions about the scope of moral standing, that is, about the boundaries of what should count in our moral deliberations.

1 } A Moral Theory Primer

In 1998, Dr. Jack Kevorkian helped Thomas Youk end his life by giving him a lethal injection of drugs—an incident that was videotaped and later broadcast on CBS’s *60 Minutes*.¹ Youk had been suffering from amyotrophic lateral sclerosis (often called Lou Gehrig’s disease), a progressive neurodegenerative disease that attacks nerve cells in the brain and spinal cord, eventually leading to death. In the later stages of the disease, its victims are completely paralyzed, as was Youk at the time of his death.

Kevorkian’s killing Youk was a case of euthanasia, which is defined as the act of killing (or allowing to die) on grounds of mercy for the victim. In this case, because Youk consented to his own death and because Kevorkian brought about Youk’s death by an act of lethal injection, Kevorkian’s action was an instance of voluntary active euthanasia. Kevorkian was eventually tried and convicted of second degree murder for his active role in bringing about Youk’s death. But even if Kevorkian did violate the law, was his action morally wrong? Youk’s immediate family and many others saw nothing morally wrong with Youk’s decision or with Kevorkian’s act. They argued, for example, that proper respect for an individual’s freedom of choice means that people in Youk’s situation have a moral right to choose to die and that, therefore, Kevorkian was not acting immorally in helping Youk end his life. Of course, many others disagreed, arguing, for example, that euthanasia is morally wrong because of its possible bad effects over time on society, including the possibility that the practice of euthanasia could be abused, and vulnerable persons might be put to death without their consent. Which side of this moral dispute is correct? Is euthanasia at least sometimes morally right, or is this practice morally wrong?

Disputes over moral issues are a fact of our social lives. Most people, through television, the Internet, magazines, and conversing with others, are familiar with some of the general contours of such disputes—disputes, for example, over the death penalty, the ethical treatment of animals, human cloning, abortion. The same sort of moral question raised about the actions of Kevorkian can be raised about these and other moral issues. Thinking critically about such moral issues is where philosophy becomes especially important.

A *philosophical* approach to moral issues has as its guiding aim arriving at correct or justified answers to questions about the morality of the death penalty, the ethical treatment of animals, human cloning, abortion, and other issues of moral concern. Given the contested nature of such practices as cloning and abortion, one needs to be able to defend one’s position with *reasons*. Just as those who dispute questions about, say, science or history are expected to give reasons for the scientific and historical beliefs they hold, those who seriously dispute moral questions are expected to give reasons for whatever moral position they

take on a certain issue. If we examine how philosophers go about providing reasons for the moral positions they take on certain issues, we find that very often they appeal to a **moral theory**. That is, in arguing for a particular position on the topic of, say, euthanasia, philosophers often make their case by applying a moral theory to the practice of euthanasia. Applying moral theory to issues of practical concern—practical issues—is one dominant way in which reasoning in ethics proceeds, and this way of tackling moral issues by applying theory to cases is featured in this book of readings.

But what is a moral theory? What are its guiding aims? What moral theories are there? How is a moral theory used in reasoning about disputed moral issues? These are the main questions of concern in this moral theory primer.

1. WHAT IS A MORAL THEORY?

According to philosopher John Rawls, “The two main concepts of ethics are those of the right and the good. . . . The structure of an ethical theory is, then, largely determined by how it defines and connects these two basic notions.”²

In explaining what a moral theory is, then, the place to begin is by clarifying the two main concepts featured in such a theory.

The Main Concepts: The Right and the Good

In ethics, the terms “right” and “wrong” are used primarily to evaluate the morality of actions, and in this chapter we are mainly concerned with moral theories that address the nature of right and wrong action (or right action, for short). Here, talk of right action in contrast to wrong action involves using the term “right” broadly to refer to actions that aren’t wrong. Used in this broad sense, to say of an action that it is right is to say that it is “all right” (not wrong) to perform, and we leave open the question of whether the act, in addition to being all right, is an action that we morally ought to perform—an obligation or duty. But we sometimes find “right” being used narrowly to refer to actions that are “the” morally right action for one to perform, and when so used, it refers to actions that are morally required or obligatory (one’s obligation or duty). Actions that are all right to perform (right in the sense of merely being not wrong) and that are also not one’s moral obligation to perform—actions that are all right to perform and all right not to perform—are morally optional. So, we have three basic categories of moral evaluation into which an action may fall: an action may be morally obligatory (something one morally ought to do, is morally required to do, is one’s duty), or morally optional, or morally wrong. To help keep this terminology straight, I have summarized what I have been saying in Figure 1.1.

Again, in ethics, the terms “good” and “bad” are used primarily in assessing the value of persons (their character) as well as experiences, things, and states of affairs. Philosophers distinguish between something’s having **intrinsic value** (that is, being intrinsically good or bad) and something’s having **extrinsic value** (that is, being extrinsically good or bad). Something has intrinsic value when its value depends on features that are *inherent* to it, whereas

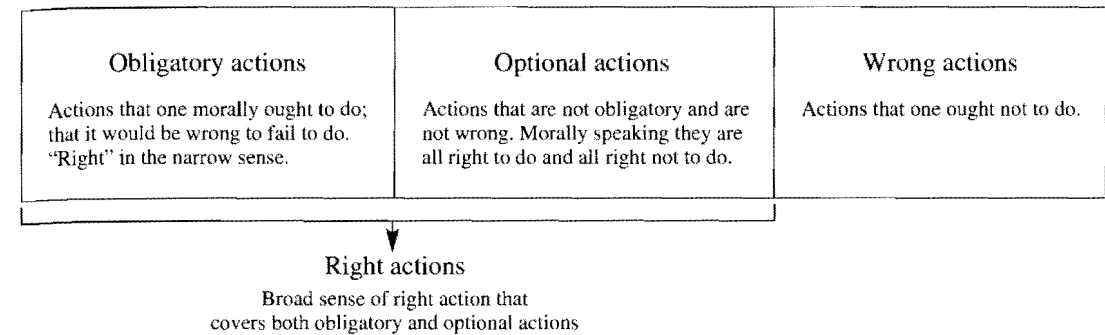


FIGURE 1.1 Basic Categories of Right Conduct

something is extrinsically good when its goodness is a matter of how it is related to something else that is intrinsically good. For instance, some philosophers maintain that happiness is intrinsically good—its goodness depends on the inherent nature of happiness—and that things like money and power, while not intrinsically good, are nevertheless extrinsically good (valuable) because they can be used to bring about or contribute to happiness. Thus, the notion of intrinsic value is the more basic of the two notions, and so philosophical accounts of value are concerned with the nature of intrinsic value. And here we can recognize three basic value categories: the *intrinsically good*, the *intrinsically bad* (also referred to as the *intrinsically evil*), and what we may call the *intrinsically value-neutral*—that is, the category of all those things that are neither intrinsically good nor bad (though they may have extrinsic value).³

A moral theory, then, is a theory about the nature of the right and the good and about the proper method for making correct or justified moral decisions. Accordingly, here are some of the main questions that a moral theory attempts to answer:

1. What *makes* an action right or wrong—what *best explains why* right acts are right and wrong acts are wrong?
2. What *makes* something good or bad—what *best explains why* intrinsically good things are intrinsically good (and similarly for things that are intrinsically bad or evil)?
3. What is the *proper method* (supposing there is one) for reasoning our way to correct or justified moral conclusions about the rightness and wrongness of actions and the goodness and badness of persons, and other items of moral evaluation?

In order to understand more fully what a moral theory is and how it attempts to answer these questions, let us relate what has just been said to the two guiding aims of moral theory.

Two Main Aims of a Moral Theory

Corresponding to the first two questions about the nature of the right and the good is what we may call the theoretical aim of a moral theory:

The **theoretical aim** of a moral theory is to discover those underlying features of actions, persons, and other items of moral evaluation that *make* them right or wrong,

good or bad and thus *explain why* such items have the moral properties they have. Features of this sort serve as *moral criteria* of the right and the good.

Our third main question about proper methodology in ethics is the basis for the practical aim of a moral theory:

The **practical aim** of a moral theory is to offer *practical guidance* for how we might arrive at correct or justified moral verdicts about matters of moral concern—verdicts which we can then use to help guide choice.

Given these aims, we can evaluate a moral theory by seeing how well it satisfies them. We will return to the issue of evaluating moral theories in section 3. For the time being, we can gain a clearer understanding of these aims by considering the role that principles typically play in moral theories.

The Role of Moral Principles

In attempting to satisfy these two aims, philosophers typically propose **moral principles**—very general moral statements that specify conditions under which an action is right (or wrong) and something is intrinsically good (or bad). Principles that state conditions for an action's being right (or wrong) are **principles of right conduct**, and those that specify conditions under which something has intrinsic value are **principles of value**. Here is an example of a principle of right conduct (where "right" is being used in its broad sense to mean "not wrong"):

P An action is right if and only if (and because) it would, if performed, likely bring about at least as much overall happiness as would any available alternative action.⁴

This principle, understood as a moral criterion of right action, purports to reveal the underlying nature of right action—what *makes* a right action right. According to P, facts about how much overall happiness an action would bring about were it to be performed are what determine whether it is morally right. Although P addresses the rightness of actions, it has implications for wrongness as well. From P, together with the definitional claim that if an action is not morally right (in the broad sense of the term) then it is morally wrong, we may infer the following:

P* An action is wrong if and only if (and because) it would, if performed, likely not bring about at least as much overall happiness as would some available alternative action.

Since, as we have just seen, principles about moral wrongness can be derived from principles of rightness, I shall, in explaining a moral theory's account of right and wrong, simply formulate a theory's principles (there may be more than one) for right action.

In addition to serving as moral criteria, principles like P are typically intended to provide some practical guidance for coming to correct or justified moral verdicts about particular issues, thus addressing the practical aim of moral theory. The idea is that if P is a correct moral principle, then we should be able to use it to guide our moral deliberations in coming to correct conclusions about the rightness of actions, thus serving as a basis for moral deci-

sion making. In reasoning our way to moral conclusions about what to do, P has us focus on the consequences of actions and instructs us to consider in particular how much overall happiness actions would likely bring about.

To sum up, a moral theory can be understood as setting forth moral principles of right conduct and value that are supposed to explain what makes an action or other object of evaluation right or wrong, good or bad (thus satisfying the theoretical aim), as well as principles that can be used to guide moral thought in arriving at correct or justified decisions about what to do (thus satisfying the practical aim).

The Structure of a Moral Theory

Finally, what Rawls calls the "structure" of a moral theory is a matter of how a theory connects the right and the good. As we shall see, some theories take the concept of the good to be more basic than the concept of the right and thus define or characterize the rightness of actions in terms of considerations of intrinsic goodness. Call such theories value-based moral theories. **Value-based moral theories** include versions of consequentialism, natural law theory, and virtue ethics. However, some moral theories do not define rightness in terms of goodness. Some theories are **duty-based moral theories**—theories that take the concept of duty to be basic and so define or characterize the rightness of actions independently of considerations of goodness. These theories are often called "deontological" moral theories (from *deon*, the Greek term for duty). The moral theory of Immanuel Kant (see later in this chapter) and theories inspired by Kant (Kantian moral theories) are arguably deontological. And what is called the ethics of prima facie duty, if not a pure deontological theory, contains deontological elements, as we shall see when we discuss this theory later in section 2.

Brief Summary

Now that we have reviewed a few basic elements of moral theory, let us briefly sum up.

- *Main concepts of moral theory.* The two main concepts featured in moral theory are the concepts of the right (and wrong) and the good (and bad).
- *Two aims of moral theory.* A moral theory can be understood as having two central aims. The theoretical aim is to explain the underlying nature of the right and the good—specifying those features of actions or other items of evaluation that *make* an action or whatever right or wrong, good or bad. We call such features "criteria." The practical aim is to offer practical guidance for how we might arrive at correct or justified moral verdicts about matters of moral concern.
- *The role of moral principles.* A moral theory is typically composed of moral principles (sometimes a single, fundamental principle) that are intended to serve as criteria of the right and the good (thus satisfying the theoretical aim) and are also intended to be useful in guiding moral thinking toward correct, or at least justified conclusions about some moral issue.
- *The structure of a moral theory.* Considerations of structure concern how a moral theory connects the concepts of the right and the good. Value-based theories make the good (intrinsic value) more basic than the right and define or characterize the right in terms of the good. Duty-based theories characterize the right independently of considerations of value.⁵

In the next section, we briefly examine six moral theories that play a large role in philosophical discussions of disputed moral issues. After presenting these theories, I devote the remaining section and an appendix to questions that are likely to occur to readers. First, there is the question of why studying moral theories is helpful in thinking about disputed moral issues when there is no *one* moral theory that is accepted by all those who study moral theory. Rather, we find a variety of apparently competing moral theories that sometimes yield conflicting moral verdicts about the same issue. So, how can appealing to moral theory really help in trying to think productively about moral issues? This is a fair question that I address in section 3. In the appendix, I briefly present two moral theories whose guiding ideas will be familiar to most all readers—the divine command theory and ethical relativism—and I explain why they aren't featured in this book's readings. However, before going on, let me say something about how one might use this chapter in studying the moral issues featured in this book.

User's Guide Interlude

In the "User's Guide," I suggested that although this chapter can be read straight through, readers may want to stop here and go on to one of the following chapters and begin their study of disputed moral issues. In the chapter introductions and the brief article summaries that precede each reading selection, I prompt readers to read (or reread) my presentations of one or more of the six moral theories I describe in the next section of this chapter. As I explained in the user's guide, I like to teach moral theory along with the readings. Seeing how a moral theory applies to a particular moral issue is helpful for understanding an author's position on the issue, which in turn helps readers gain a deeper understanding of and an appreciation for moral theory.

As for integrating section 3 and the appendix, I recommend consulting these parts of the chapter when the questions they address are prompted by one's thinking about and discussing the book's readings.

2. SIX ESSENTIAL MORAL THEORIES

Six types of moral theory are prominently represented in our readings: consequentialism, natural law theory, Kantian moral theory, rights-based moral theory, virtue ethics, and the ethics of *prima facie* duty. Here, then, is an overview of these various theories that will provide useful background for understanding our readings.

A. Consequentialism

In thinking about moral issues, one obvious thing to do is to consider the consequences or effects of various actions—the consequences or effects on matters that are of concern to us. **Consequentialism** is a type of moral theory according to which consequences of actions are all that matter in determining the rightness and wrongness of actions. Its guiding idea is this:

- C Right action is to be understood entirely in terms of the overall intrinsic value of the consequences of the action compared to the overall intrinsic value of the consequences associated with alternative actions an agent might perform instead. An action is right if and only if (and because) its consequences would be at least as good as the consequences of any alternative action that the agent might instead perform.

A number of important ideas are packed into C that we need to unpack—ideas that are present in the varieties of consequentialist moral theory presented next. Let us sort them out.

- First, consequentialist moral theory is a *value-based moral theory*: it characterizes or defines right action in terms of intrinsic value.
- Second, this sort of theory involves the fairly intuitive idea of *alternative actions* open to an agent: in circumstances calling for a moral choice, an agent is confronted by a range of alternative actions, any one of which she might choose to perform.
- Third (and relatedly), consequentialism is a *comparative* theory of right action: the rightness (or wrongness) of an action depends on how much intrinsic value it would likely produce (if performed) compared to how much intrinsic value alternative actions would likely produce (if performed).
- Fourth, the consequentialist account of right action is a *maximizing* conception: we are to perform that action, from among the alternatives, whose consequences will have *at least as much* overall value as any other.
- Fifth, and finally, consequentialism is a strongly *impartialist* moral theory in the sense that the rightness or wrongness of an action is made to depend on the values of the consequences for *everyone* who is affected by the action, where everyone affected counts *equally*. (This fifth point will become clearer when we consider particular versions of consequentialism.)

Consequentialism, we have noted, is a *general type* of moral theory that has a variety of species. For instance, consequentialists may differ over the issue of what has intrinsic value. Those versions that take happiness or welfare alone to have intrinsic value are versions of utilitarianism, whereas those that take human perfection to have intrinsic value are versions of perfectionism. Again, consequentialists may differ over the primary focus of consequentialist evaluation. Some versions focus on individual actions, other versions focus on rules. So, we can distinguish four main species of consequentialism. Let us explore further.

Utilitarianism has been perhaps the most prominent form of consequentialism, so let us begin with it.

Utilitarianism

Utilitarianism was originally developed and defended by Jeremy Bentham (1748–1832) and later refined by John Stuart Mill (1806–1873).⁶ Their basic idea is that it is *human welfare* or *happiness* that alone is intrinsically valuable and that the rightness or wrongness of actions depends entirely on how they affect human welfare or happiness. As a consequentialist theory, utilitarianism requires that one *maximize* welfare where the welfare of *all* individuals who will be affected by some action counts. We can sharpen our characterization of this theory by introducing the technical term "utility," which refers to the *net value* of the consequences of actions—how much overall welfare or happiness would likely result from an

action, taking into account both the short-term and long-term effects of the action on the welfare of all who will be affected. The basic idea is that the moral status of an action—its rightness or wrongness—depends both on how much happiness (if any) it would likely produce for each individual affected were it to be performed, as well as how much unhappiness (if any) it would likely produce for each affected person were it to be performed. For each alternative action, then, we can consider the *net balance* of overall happiness versus unhappiness associated with that action. Call this overall net value the **utility** of an action. We can now formulate a generic statement of the basic utilitarian principle—the **principle of utility**:

- U** An action is right if and only if (and because) it would (if performed) likely produce at least as high a utility (net overall balance of welfare) as would any other alternative action one might perform instead.⁷

Notice that the utility of an action might be negative. That is, all things considered, an action may produce a net balance of unhappiness over happiness were it to be performed. Moreover, since U (like all versions of C) is comparative, it may turn out that the right action in some unfortunate circumstance is the one that would likely bring about the least amount of overall negative utility.

As formulated, U leaves open questions about the nature of happiness and unhappiness about which there are different philosophical theories.⁸ Bentham and (apparently) Mill held that happiness is entirely constituted by experiences of pleasure and unhappiness by experiences of displeasure or pain. And so their theory of intrinsic value is called **value hedonism**: *only* states of pleasure have positive intrinsic value and *only* states of pain have intrinsic negative value; anything else of value is of mere extrinsic value. So, for instance, for the value hedonist, any positive value that knowledge may have is extrinsic: it is only of positive value when it contributes to bringing about what has intrinsic value, namely pleasure (or the alleviation of pain). It should be noted that a value hedonist need not (and should not) take an excessively narrow view of pleasure and pain; the hedonist can follow Bentham and Mill in holding that in addition to such bodily pleasures of the sort one gets from eating delicious food or having a massage, there are aesthetic and intellectual pleasures such as appreciating a beautifully written poem. Moreover, the value hedonist will recognize not only passive pleasures of the sort just mentioned, but also active pleasures as when one plays a game or is involved in some creative activity. So value hedonism can recognize a broad range of pleasurable experiences that have positive intrinsic value and a broad range of painful experiences that have negative intrinsic value.

If we now combine the principle of utility (U) with value hedonism, we obtain **hedonistic utilitarianism**:

- HU** An action is right if and only if (and because) it would likely produce (if performed) at least as high a net balance of pleasure (or less pain) as would any other alternative action one might do instead.

But as I hope my presentation has made clear, one need not accept hedonism as a theory of value in order to be a utilitarian. In fact, many contemporary utilitarians reject value hedonism and accept some other conception of happiness or welfare. But, again, what makes a theory a version of utilitarianism is that the theory accepts the basic consequentialist claim, C, together with the idea that it is human happiness or human well-being that has intrinsic value and is to be promoted in what we do.

Perfectionist Consequentialism

But a consequentialist need not be a utilitarian—she might hold that there are items having intrinsic value other than happiness that are important in determining the rightness or wrongness of action. To illustrate, I have chosen what is called **perfectionist consequentialism**—a species of the generic view that accepts a perfectionist theory of value.⁹ According to a **value perfectionist**, it is states of human perfection, including knowledge and achievement that have intrinsic value.¹⁰ One might come to have a great deal of knowledge and achievement in one's life, yet not be happy. So a perfectionist theory of the good is not the same as a happiness theory of the good. We might formulate the basic principle of perfectionist consequentialism as follows:

- PC** An action is right if and only if (and because) it would (if performed) likely bring about a greater net balance of perfectionist goods than would any alternative action one might perform instead.

The distinction between utilitarianism and perfectionist consequentialism has to do with differences over what has intrinsic value for purposes of morally evaluating actions. And notice that the consequentialist principles presented thus far refer to particular concrete actions and their consequences, so the views (expressed in principles U, HU, and PC) are versions of **act consequentialism**. However, as mentioned at the outset, another important division within the ranks of consequentialists is between act and rule versions of the view. So let us turn from act versions to rule versions.

Rule Consequentialism

Moral rules—rules, for example, against lying, theft, and killing—are generally thought to be significant in thinking about particular moral issues. The importance of moral rules is emphasized by rule consequentialists. Whereas act consequentialism is the view that the rightness of a particular, concrete action—an actual or possible doing by a person at a time—depends on the value of its consequences, **rule consequentialism** is the view that the rightness or wrongness of an action depends on whether it is required, permitted, or prohibited by a rule whose consequences are best.¹¹ So rule consequentialism involves two levels of evaluation: first, rules that require, permit, or prohibit various courses of action are evaluated by reference to the values of their consequences, and second, a particular action is evaluated by determining whether it is required, permitted, or prohibited by a rule whose consequences are best. Let us explore this view a bit further.

The sense in which a rule can have consequences has to do with the fact that were people to accept the rule in question, this would influence what they do. So, we can evaluate a rule by asking what consequences would likely be brought about were it to be generally accepted in society. Call the value associated with rules their **acceptance value**. This idea is familiar. Think of debates in the sporting world about changing the rules of some sport. The focus in such debates is on the likely effects the proposed rule change would have on the game, were it to be accepted.

According to rule consequentialism, then, the morality of a particular action in some situation depends upon the acceptance values of various competing rules that are relevant to the situation in question. We can thus formulate this theory with the following principle of right conduct:

RC An action is right if and only if (and because) it is permitted by a rule whose associated acceptance value is at least as high as the acceptance value of any other rule applying to the situation.

In order to better understand this principle, let us illustrate its application with a simple example.

Suppose that I have promised to help you move next Friday morning. Friday morning arrives, and many alternative courses of action are open to me. Among them are these:

- A1. Keep my promise (and show up at your place),
- A2. Break my promise (and do something else).

Corresponding to each of these alternative actions, we have these rules:

- R1. Whenever one makes a promise, keep it,
- R2. Whenever one makes a promise, break it if one feels like it.

Now consider the acceptance values associated with these rules. I think we can all agree that acceptance value of **R1** is far greater than that of **R2**. So (ignoring for the moment that there may be other competing rules to be considered in this situation) rule consequentialism implies that one ought to keep one's promise.

Finally notice that act and rule consequentialism may diverge in their moral implications. To stick with the previous example, suppose that by breaking my promise and instead hanging out with some friends at the local pool hall will likely produce a greater level of overall intrinsic value than would the backbreaking work of helping you move. Besides, you've lined up plenty of help; I won't be missed that much. Act consequentialism implies that it would be morally permissible to go ahead and break the promise; rule consequentialism by contrast implies that I am morally obliged to keep my promise.

Brief Summary

Let us pause for a moment to summarize (see Fig. 1.2) what we have covered. As we have seen, the basic consequentialist idea (C) can be developed in a variety of ways; we have considered four versions of this generic approach to ethics.

For now, the main idea to take away from this discussion is that for all varieties of consequentialism, the rightness or wrongness of an action depends entirely on the net intrinsic value of the consequences of either individual actions or rules. Consequentialist theories (and especially utilitarianism) are often discussed in articles and books about disputed moral issues. Some authors appeal to consequentialism to justify their particular views on some moral issue; other authors will contrast their approach with consequentialism.

Applying Consequentialism

To convey a sense of how one is to go about applying consequentialism to a particular moral issue, let us work with act utilitarianism as expressed earlier in U. And to make things fairly manageable, let us consider a rather simple case.

Suppose that I am in charge of inviting a guest philosopher to speak at my university, and that I've narrowed the choices to two. On the one hand, I can invite Dr. Brilliant, a very well-




FIGURE 1.2 Some Forms of Consequentialism

known and innovative philosopher but whose manner of presentation is decidedly dull. The philosophy faculty will no doubt take pleasure in his presentation and will benefit intellectually from what he has to say, but others will be bored stiff and get little out of the talk. On the other hand, I can invite Dr. Flash who is not nearly as accomplished as Dr. Brilliant but who I know is an extremely engaging speaker. Suppose that five professional philosophers and forty-five students are expected to attend the lecture no matter which of these two philosophers I invite.

Now if I apply U to my situation, it would have me invite the speaker whose talk will produce the greatest amount of overall happiness. A careful application of U would require that I consider each person who will be affected by the lecture of Dr. Brilliant and determine how much happiness (if any) that person would experience as a result of this lecture and then determine how much unhappiness (if any) that person would experience as a result of this lecture. Once I have done this calculation for each person, I then calculate how much total happiness the lecture would cause and how much total unhappiness it would cause in order to arrive at the overall net value associated with Dr. Brilliant's lecture. I do the same for Dr. Flash. The lecture I ought to sponsor and hence the philosopher I ought to invite depends on which talk will result in the greatest amount of intrinsic value.

Obviously, to apply U to one's own choices with any precision would require much factual information: information about one's circumstances, about the various alternative actions one might perform in one's particular circumstances, about the individuals who will be affected either negatively or positively were one to perform a particular action, and about the overall amount of happiness or unhappiness, both short term and long term, that would likely result from each of the various alternative actions. Some critics of consequentialism argue that when it comes to satisfying the practical aim of moral theory—the aim of providing practical guidance for arriving at correct or justified verdicts for the further aim of acting on such verdicts—consequentialism makes implausible demands on what one needs to know in order to apply it to particular cases. Defenders reply that even if *precise* application of this sort of moral theory is not feasible, we can and do make rough estimates of the values of the consequences of various alternative actions, and that in doing so we must realize that the moral verdicts we reach as a result are likely to be highly fallible. But this, say the defenders, is something we just have to live with given our limited information about the effects of our actions.

B. Natural Law Theory

The idea that certain actions or practices are “natural” while others are “unnatural” is commonly offered as a reason why certain “unnatural” actions are wrong and that we ought to do what is natural. Think of popular arguments against homosexuality. This idea of morality being natural is associated with the **natural law theory**.¹²

This type of moral theory is often traced to the thirteenth-century philosopher and theologian St. Thomas Aquinas (1225–1274). It gets its name from the guiding idea that there are objectively true moral principles that are grounded in human nature.¹³ Because there are objective facts about human nature that determine what our good consists in, and because moral requirements have to do with maintaining and promoting the human goods, these requirements, unlike the rules of some club or made-up game, are part of the natural order. Because the natural law theory bases right action on considerations of intrinsic value, it is a value-based theory of right conduct, as is consequentialism. However, as we shall see in setting out this theory, natural law theory is opposed to consequentialism—it denies that the *only* considerations that matter when it comes to right action are consequences. So, to understand this theory let us proceed by first presenting its theory of intrinsic value and then presenting its theory of right conduct in two parts: (a) first, the “core” of theory and then (b) the doctrine of double effect.

*Theory of Intrinsic Value*¹⁴

According to Aquinas’s version of natural law theory, there are four basic intrinsic goods:

- Human life
- Human procreation (which includes raising children)
- Human knowledge
- Human sociability (this value has to do with associations and bonds with others, including friendship, social organizations, and political organizations)

Each of these items, then, has intrinsic value and their destruction is intrinsically bad or evil. These four values are the basis for the core of natural law theory.

The Core

We can state the basic principle of natural law theory roughly as follows:

NLT An action is right if and only if (and because) in performing the action one does not directly violate any of the basic values.

Thus, killing a human being (with some exceptions explained later) is morally wrong. If we suppose, as many natural law theorists do, that the use of contraceptives thwarts human procreation, then their use is morally wrong. Interfering with the good of knowledge by distorting information or by lying is morally wrong. Destroying legitimate social bonds through the advocacy of anarchy is morally wrong.

But what about hard cases in which no matter what one does, one will violate at least one of the basic values and thus bring about evil through whichever action one chooses? Let us consider a much discussed case involving abortion. Suppose that a pregnant woman has cancer of the uterus and must have a hysterectomy (removal of her uterus) to save her life.

Human life is one of the intrinsic goods, so having the operation will have at least one good effect. But suppose (just for the sake of the example) that from conception the fetus counts as a human life and so having the hysterectomy would bring about the death of the unborn human life. This effect, because it involves the destruction of something having intrinsic value—human life—is an evil. And let us suppose that this moral dilemma is unavoidable in this case because there is no other way to save the woman’s life while also preserving the life of her fetus. How does the natural law theory deal with this kind of case? After all, the core of the theory seems to say that any action that violates one or more of the basic goods is wrong, period. But if it really does say this, then we have to conclude that her having the operation is wrong, but also her not having the operation is wrong (because she will fail to preserve her own life). How can natural law theory deal with this moral dilemma?

If we go back and inspect the basic principle of natural law theory, NLT, we notice that what it prohibits are actions that *directly* violate one or more of the basic goods, thereby bringing about evil. But what counts as a direct violation? Can there be morally permissible “indirect” violations? These questions bring us to the next major component of natural law ethics—the doctrine of double effect.

The Doctrine of Double Effect

In addition to the core principle (NLT), the natural law theory also embraces the following set of provisions that compose the **doctrine of double effect**—so named because it concerns cases in which performing an action would have at least one good effect and one bad effect (where good and bad have to do with the theory’s list of intrinsic goods). So this doctrine is meant to address the question of whether it is ever morally permissible to knowingly bring about bad or evil consequences where one’s aim in action is to bring about or preserve one or more of the basic human goods. Here, then, is a statement of the various provisions making up the doctrine:

DDE An action that would bring about at least one evil effect and at least one good effect is morally permissible if (and only if) the following conditions are satisfied:

Intrinsic permissibility: The action in question, apart from its effects, is morally permissible;

Necessity: It is not possible to bring about the good effect except by performing an action that will bring about the evil effect in question;

Nonintentionality: The evil effect is not intended—it is neither one’s end nor a chosen means for bringing about some intended end;

Proportionality: The evil that will be brought about by the action is not out of proportion to the good being aimed at.

What this principle does is help define the idea of a direct violation of a human good which is the central idea in the core principle, NLT. We shall return to this point in a moment. For the time being, let us explain DDE by showing how it would apply to the case just described.

In applying DDE to our moral dilemma, we must ask whether all four of the doctrine’s provisions are satisfied. Let us take them in order. (1) First, since having a hysterectomy is not an intrinsically wrong action, the first requirement is satisfied. (2) Furthermore, given my description of the case, the second requirement of DDE is met because having a hysterec-

tomy is the *only* way to save her life. Were there some other operation or some medication that would both save the woman's life and preserve the life of the fetus, then the necessity condition would not be met and the hysterectomy would be wrong. But we are supposing that there are no such options in this case.

(3) The third requirement rests on the distinction between effects that one intends in action and effects that may be foreseen but are unintended. One intends some consequence or effect when either it is something one is aiming to bring about (an end) or it is one's chosen means for bringing about some desired end. Here is a simple, everyday example. I fire a rifle in order to hit the paper target, but in so doing I know that the noise from the rifle will frighten nearby wildlife. But even though I can foresee that my act of pulling the trigger will frighten those animals, this effect is not intended: it is not my purpose—my purpose is to hit the target, and their being frightened is not a means for achieving my end—the means is taking aim and firing. So the effect of my act of firing—frightening those animals—is not something I intend, rather it is a foreseen but unintended side effect of what I do.

Returning now to our example, we find that this third provision is satisfied because although the death of the unborn child is a foreseen effect of the hysterectomy, its death is not her chief aim or end (saving her own life), and it is not a means by which her life will be saved. After all, were she not pregnant, she would still have the operation to save her life, and so the death of the unborn is a mere unintended and unfortunate side effect of the operation. Removing the cancer is what will save her life.

(4) Finally, the evil that will result from the operation (loss of one innocent human life) is not grossly out of proportion to the good that will result (saving an innocent human life). (When DDE is applied to the morality of war activities, considerations of proportionality of evil to good become especially relevant. See the introduction to chapter 10 and the articles in that chapter by Haig Khatchadourian and James P. Sterba.)

Having explained the DDE, we can now return to the core principle, NLT, and explain how these two elements are related in natural law ethics. The idea is that, according to NLT, we are not to *directly* violate any of the basic human goods. The DDE helps define what counts as a direct violation: direct violations are those that cannot be justified by the doctrine of double effect.

Before going on, it will be useful to pause for a moment to compare the natural law theory with consequentialism. In response to our moral dilemma involving the hysterectomy, an act consequentialist will say that we should consider the value of the consequences of the alternative actions (having a hysterectomy or refraining from this operation) and choose the action with the best consequences. In short, for the act consequentialist good results justify the means. But not for the natural law theorist, because on her theory one may not act in direct violation of the basic goods even if by doing so one would produce better consequences. Good ends do not always justify an action that is a means to those ends. For instance, I am not permitted to intentionally kill one innocent human being (do evil) even if by doing so I can save five others (bring about good). To see how consequentialism and natural law theory yield different verdicts about a difficult moral case, consider the case of a woman who is pregnant, but this time she is suffering from a "tubal" pregnancy, which means that her fetus is lodged in her fallopian tube and thus has not implanted itself into the uterine wall. If nothing is done, both fetus and woman will die. The only thing that can be done to save the woman is to remove the fetus, which will bring about its death. Exercise: apply act consequentialism and the natural law theory to this case to see whether they differ in their moral implications.

Applying Natural Law Theory

In applying the natural law theory to some case in order to determine whether a particular course of action is morally right, one begins with the core principle, NLT, and asks whether the action in question would violate any of the basic goods. If not, then the action is not wrong. But if it would violate one or more of the basic goods, then one has to determine whether the action would constitute a *direct* violation. And to do that, one makes use of the DDE. If the action satisfies all four provisions of DDE, then the violation is not direct and the action is morally permissible. If the action does not pass DDE, then the action involves a direct violation of one or more of the intrinsic goods and is, therefore, wrong.

Of course, as with all moral theories, applying the natural law theory is not a mechanical process. For one thing, one must address questions about the proper interpretation of the four basic human goods. Surely coming to have knowledge of the basic laws that govern the physical universe is intrinsically valuable, if any knowledge is. But what if, for example, I spend my time counting the number of needles on a cactus plant for no particular reason; is the knowledge I acquire about the number of needles really of any intrinsic value? One can raise similar questions about the other three basic human goods. Furthermore, applying the doctrine of double effect raises questions of interpretation. For instance, the proportionality provision requires that the evil caused by an action not be "out of proportion" to the good effects of that action. But, of course, determining when things are out of proportion requires sensitivity to particular cases and the use of good judgment.

These points about interpretation are not meant as a criticism of natural law theory; rather they call attention to the fact that applying it to particular moral issues requires that we interpret its various elements. As we shall see, a similar point applies to Kantian moral theory.

C. Kantian Moral Theory

Most everyone has come across moral arguments that appeal to the **golden rule**: do unto others as you would have them do unto you. This rule encapsulates a kind of test for thinking about the morality of actions: it asks the individual making a moral choice that will affect others to consider how one would like it were one on the receiving end of the action in question. In the case of Thomas Youk with which we began the chapter, the golden rule would have Kevorkian consider what he would want done to (or for) him were he in Youk's situation. Various objections have been made to the golden rule—for instance, it suggests that the rightness or wrongness of an action depends simply on what one does or would desire. But people can have crazy desires. A masochist who inflicts pain on others might cheerfully say that he would have others do unto him as he is doing to them. Do we want to conclude that his causing others pain is morally right? Perhaps there is some interpretation of the golden rule that does not yield the wrong result in the case of the masochist or other examples that have been used against it. Nevertheless, there is something about the *spirit* of the golden rule that seems right. The idea suggested by this rule is that morality requires that we not treat people unfairly, that we respect other persons by taking them into account in our moral deliberations. This suggestion is quite vague but finds one articulation in Kantian moral theory to which we now turn.

Kantian moral theory derives from the moral writings of the German philosopher Immanuel Kant (1724–1804), which continue to have an enormous influence on contempora-

rary ethics.¹⁵ Central to Kant's moral theory is the idea that moral requirements can be expressed as commands or imperatives that categorically bid us to perform certain actions—requirements that apply to us regardless of what we might happen to want or desire or how such actions bear on the production of our own happiness. Kant thought that specific moral requirements could be derived from a fundamental moral principle that he called the **categorical imperative**. Moreover, Kant offered various alternative formulations of his fundamental moral principle. The two I will consider are the ones that are most often applied to moral issues.

The Humanity Formulation

One of Kant's formulations of his categorical imperative is called the **Humanity formulation**:

H An action is right if and only if (and because) the action treats persons (including oneself) as ends in themselves and not as mere means.

Obviously, to make use of this principle, we need to know what it means to treat someone as an end and what it means to treat someone as a mere means. Space does not permit a thorough discussion of these ideas, so a few illustrations will have to suffice.¹⁶

Deception and coercion are two ways in which one can treat another person as a mere means—as an object to be manipulated. Deceiving someone in order to get him or her to do something he or she would otherwise not agree to do constitutes using the person as though that person were a mere instrument at one's disposal for promoting one's own ends. Again, many cases of coercing someone by threats involves attempting to manipulate that person for one's own purposes and hence constitutes an attempt to use him or her as a mere means to one's own ends.

But Kant's Humanity formulation requires not only that we not treat others as mere means to our own ends (a negative requirement), but also that we treat them as ends in themselves (a positive requirement). For Kant, to say that persons are ends in themselves is to say that they have a special worth or value that demands of us that we have a certain positive regard for them. Kant refers to this special worth as *dignity*.¹⁷ So, for instance, if I fail to help those who need and deserve my help, I don't treat them as mere means, but I do fail to have a positive regard for their welfare and thus fail to properly recognize their worth as persons.

Applying Kant's Humanity Formulation

As just explained, applying the Humanity formulation requires consideration of the dual requirements that we not treat people as mere means and that we also not fail to treat them as ends in themselves—as individuals who have dignity. Interpreting these requirements is where the hard work comes in: what are the boundaries when it comes to treating people as *mere* means? If, in walking by, I see that you are wearing a watch and ask you for the time, I am using you as a means for finding out what time it is, but I am not thereby using you as a *mere* means to my own ends. We have noted that deception and coercion represent main ways in which one might use someone as a mere means—as something to be manipulated. So we have a good start on interpreting the idea of treatment as mere means. Here is not the place to consider other ways in which our actions might involve treating someone as a mere means. Rather, the point I wish to make is that we have some idea of what it means to treat

someone as a mere means and we must build on this understanding to apply the Humanity formulation to a range of moral issues.

Similar remarks apply to the requirement that we positively treat others as ends in themselves. Here it is interesting to note that Kant argued that in satisfying this requirement, we are obligated to adopt two very general goals—the goal of promoting the (morally permissible) ends of others, and the goal of self-perfection. Such wide-open goals allow a person much latitude in deciding in what ways and on what occasions to promote the ends of others and one's own self-perfection. For Kant, then, applying the positive requirement embedded in H is a matter of figuring out how best to integrate the promotion of the well-being of others and one's own self-perfection into a moral life.

The Universal Law Formulation

Kant's other main formulation of the categorical imperative, the **Universal Law formulation**, expresses a test whereby we can determine whether our actions are right or wrong.

UL An action is right if and only if one can both (a) consistently conceive of everyone adopting and acting on the general policy (that, is the maxim) of one's action, and also (b) consistently will that everyone act on that maxim.¹⁸

This formulation will remind readers of the golden rule, though notice that UL does not refer to an agent's wants; rather it represents a kind of consistency test.¹⁹ Unfortunately, interpreting Kant's two-part test requires some explanation. So let me say a bit more about UL and then, using some of Kant's own examples, show how it can be applied.

According to Kant, when we act, we act on a general policy that is called a "maxim." To determine the morality of an action, one formulates the general policy of one's action and asks whether one could consistently both conceive of and will that everyone act on the same policy or, to put it in Kant's terms, one asks whether one could consistently conceive and will that the maxim of one's action become a "universal law" governing everyone's behavior. If so, then the action is right; if not, then the action is wrong. So UL expresses a two-part test one can use to determine the rightness or wrongness of actions. To make Kant's tests more concrete, let us consider a few of Kant's own sample applications of UL.

One of Kant's examples involves making a lying promise—that is, a promise that one has no intention of keeping. Consider a case in which I desperately need money right away and the only way I can get it is by getting a loan which I must promise to repay. I know, however, that I won't be able to repay the loan. The maxim corresponding to the action I am considering is:

M1 Whenever I need money and can get it only by making a lying promise, I will borrow the money by making a lying promise.

Kant's principle, UL, would have me test the morality of making a lying promise by asking whether I could *consistently conceive* and *will* that everyone act on M1—that everyone who needs money in such circumstances as mine make a lying promise. Let us first ask whether this is something I can consistently conceive. If it isn't, then I certainly can't consistently will that everyone adopt and act on it.

Kant claims that when I think through what would be involved in everyone acting on M1, I realize that I cannot even consistently conceive of a world in which everyone in need of

money successfully makes lying promises. After all, a world in which everyone in need of money goes around trying to get the money by making a lying promise is one in which successful promising becomes impossible since, as Kant observes, “no one would believe what was promised him but would laugh at all such expressions as vain pretenses.”²⁰ Thus, trying to even conceive of a world in which everyone in need of money acts on M1 involves an inconsistency: it is a world in which (1) *everyone* in need gets money by making a lying promise, but because of the breakdown in the institution of promising that would result, it is a world in which (2) *not everyone* in need gets money by making a lying promise for the reason Kant gives. But if I can’t consistently conceive of everyone acting on M1, then my maxim fails the first test mentioned in UL. And if I can’t consistently conceive that everyone act on M1, this shows me that, in making a lying promise, I am acting on an immoral policy and that my action is wrong.

But why is the fact that one cannot consistently conceive that everyone act on one’s maxim an indication that the action in question is wrong? Kant’s idea here seems to be that in performing an action whose maxim I cannot consistently conceive everyone adopting, I am, in effect, proposing to make an exception of myself—an exception that I cannot justify. In making an exception of myself, I am failing to respect others because I’m taking advantage of the fact that many others do not make lying promises. And so these reflections lead us to conclude that making a lying promise is morally wrong.

Here is another example Kant uses to illustrate the application of UL that has to do with clause (b) of UL. Suppose I am in a position to help someone in need but would rather not be bothered. The maxim Kant has us consider is:

M2 Whenever I am able to help others in need, I will refrain from helping them.

Using UL, I am to consider whether I can consistently conceive of a world in which everyone adopts and acts on this maxim. Is such a world conceivable? It would seem so. Granted, a world in which people in need did not receive help from others would be a very unpleasant place. Perhaps the human race would not survive in such a world. But we can certainly conceive of a world in which the human race ceases to exist. So, M2 passes the first part of Kant’s UL test.

But can one *will* that M2 be adopted and acted upon by everyone? Upon reflection I realize that if I will that everyone adopt and act on M2, I am thereby willing that others refuse to help me when I am in need. But willing that others refuse to help me is inconsistent with the fact that as a rational agent I do will that others help me when I am in need. That is, as a rational agent, I embrace the following maxim:

RM I will that others who are able to do so help me when I am in need.

But an implication of my willing that everyone adopt and act on M2 would be:

IM I will that others who are able to help me, refuse to do so when I am in need.

RM is inconsistent with IM—and IM an implication of willing that everyone adopt and act on M2. Thus, I cannot consistently will that everyone adopt and act on M2. Since I cannot consistently will that everyone adopt M2, then according to clause (b) of Kant’s UL, my action of refusing to help others in need is morally wrong.

What is the point of Kant’s UL formulation involving two tests? Kant thought that these two tests could distinguish between what he called “narrow” or “perfect” duty and “wide” or “imperfect” duty: maxims that one cannot consistently conceive as adopted and acted on by

everyone involve actions that are contrary to narrow duty, whereas those that can be so conceived but which one cannot consistently will involve actions that are contrary to wide duty. The realm of narrow duty concerns those actions and omissions regarding which one has comparatively little room for when and how one complies with the duty. If I have promised to do something for you on a particular occasion—help you with your taxes—then to fulfill my obligation I must perform some rather specific action (helping you with your taxes) at a certain time. By contrast, a wide duty is one which can be fulfilled in a variety of ways and situations, giving one much leeway in how and when to fulfill the duty. Duties of charity—helping others—are like this.

It is important to notice that in his examples, Kant is not arguing that if everyone went around making lying promises the consequences would be bad and therefore making a lying promise is wrong. Again, he does not argue that the consequences of everyone refusing to help others in need would be bad and therefore refusing help to others is wrong. Such ways of arguing are characteristic of consequentialism, but Kant rejects consequentialism. He has us consider the implications of everyone acting on the general policy behind one’s own action because he thinks doing so is a way of revealing any inconsistencies in what one wills, which in turn indicates whether an action fails to respect persons. So the test involved in the categorical imperative is meant to reveal whether one’s action shows a proper respect for persons.

Applying the Universal Law Formulation

Since this formulation expresses tests for determining the rightness or wrongness of actions, I have been illustrating how it is to be applied in thinking through moral issues. If we step back from these illustrations, we can summarize the basic procedure to be followed. In applying UL to some actual or contemplated action of yours, here are the basic steps to follow:

- Formulate the maxim on which you are proposing to act, which will have the form “I will ____ whenever ____,” where the blanks are filled with a description of your action and circumstances, respectively.
- Next, you consider the possibility of everyone in your circumstances adopting and acting on that same maxim. In particular, you ask yourself whether you can consistently *conceive* of a world in which everyone adopts and acts on the maxim in question. This is the test expressed in clause (a) of UL.
- If you cannot even conceive of such a world, then action on the maxim is morally wrong—a violation of narrow duty. The lying promise example illustrates this result.
- If you can consistently conceive of a world in which everyone adopts and acts on the maxim, then you are to ask yourself whether you could, even so, consistently will that everyone adopt and act on the maxim. This is the test expressed in clause (b) of UL.
- If you cannot consistently will that everyone adopt and act on that maxim, then action on the maxim is wrong—a violation of wide duty. The case of refusing to help others illustrates this result.
- Finally, if your maxim is such that you can both consistently conceive of a world in which everyone adopts and acts on your maxim and consistently will this to be the case—if the maxim passes both tests—then action on the maxim is morally right.

The two main challenges for anyone applying the UL formulation to a particular issue is to correctly formulate one’s maxim and then carefully think through Kant’s two consistency tests.

How are the two formulations of the categorical imperative—H and UL—related? They are supposed to be alternative formulations of the same basic principle, rather than two entirely distinct principles. One way to see how this might be so is to notice that in cases in which I cannot consistently conceive or will that my maxim be adopted by everyone, I am making an exception of myself, and in doing so either I am treating someone as a mere means or I am failing to treat others as ends in themselves. And, of course, treating others as mere means and failing to treat them as ends in themselves is precisely what the Humanity formulation rules out.

In our readings in which authors make use of Kant's categorical imperative in arguing for some position, most of them appeal to the Humanity formulation, which has wide intuitive appeal. However, one of our authors, Harry Gensler, makes use of Kant's Universal Law formulation in making a case for the wrongness of most abortions.

D. Rights-Based Moral Theory

In our brief survey of moral theories up to this point, we have seen how those theories attempt to give accounts of the nature of right (and wrong) action, where "right" is being used in its adjectival sense. Nothing so far has been said about the notion of a *right* or *rights*. However, in moral theory we distinguish between the adjectival use of "right" as in *right action* and its use as a noun as in *a right to life*. One can hold that actions may or may not be right—actions that are either permissible or impermissible to do—without also holding that there are things called rights. Let us now focus on this latter concept by indicating what a right is supposed to be, and then we will be in a position to say something about the role of rights in moral theory, including the possibility of a rights-based moral theory.

Rights

Perhaps the most basic idea of a **right** is that of a *claim* of one person against some other party. It is common to recognize both negative and positive rights. A negative right, then, held by someone *A* to perform some action *X*, would be a claim by that person that others refrain from interfering with *A*'s performing *X*. A right to certain liberties such as free speech is an example of a negative right—a right to be free from the interference of others in expressing oneself. A positive right of *A* against *B* would involve *A*'s claim that *B* do or provide something (some service or some thing) to *A*. An example would be a right to the basic necessities (for example, food and shelter) of survival.

Two further points about rights are worth noting. First, there are philosophical disputes about the *scope* of rights holders—those creatures and possibly things that *have* rights. According to some views that would restrict the scope of rights holders, only those creatures that have a developed capacity to reason can be the holders of rights. Less restrictive views would allow that anything having interests, including certain nonhuman animals, can be rights holders. Among those who approach moral issues from the perspective of rights—including the issues of abortion, animals, and the environment—we find important differences in views about the requirements for holding rights.

Second, it is also important to distinguish legal from moral rights. Legal rights require a legal system that recognizes certain claims as having legal status; moral rights do not. The

laws of a country, for instance, may afford no legal rights to some of its citizens who nevertheless have moral rights. In debates over the morality of various activities and practices where talk of rights enter the discussion, one is mainly concerned with moral rights.

Rights and Moral Theory

A great deal of contemporary discussion about moral issues—particularly those in dispute—is couched in terms of rights. Does a human fetus have a moral right to life? Does a terminally ill patient in severe pain have a moral right to die? Do people have a moral right to reproduce by cloning? Do animals have moral rights? In this section, I want to explain how rights figure into moral theories. In doing so, I will make two main points: (1) All of the moral theories we have already surveyed (as well as the two that follow) can recognize moral rights. However, (2) what is distinctive of a *rights-based* moral theory is that it takes rights to be in some sense more basic than right action and duty.²¹ Let us take these points one by one.

First, a utilitarian who recognizes moral rights will attempt to explain rights on the basis of utility by claiming that a moral right is a kind of claim that is justified by the fact that its recognition will contribute to the maximization of overall welfare. This means that for the utilitarian—who, as a consequentialist, embraces a value-based moral theory—rights are derivative rather than basic in her moral theory. Similar remarks apply to many if not all of the moral theories featured in this chapter. This brings us to the second point—one concerning rights-based moral theories.

Rights-Based Moral Theory

Might there be a **rights-based moral theory**—a moral theory according to which rights are more basic than right action (and duty)? Unlike the other theories featured in this chapter, rights-based theories are relatively underdeveloped despite the fact that appeals to rights are very common in discussions of moral issues. What we find in the writings of authors who appeal to rights in discussing particular moral issues is that they often fail to indicate the nature of rights—whether they have a consequentialist, natural law, Kantian, or some other basis on the one hand, or whether, on the other hand, they are conceived as basic in the theory. So let us consider the idea of a rights-based moral theory.

According to such a theory, rights are more basic than right action and duty. But one might think that duties *must* be more basic than rights, and so there cannot be a rights-based moral theory. After all, as explained earlier, a moral right is a claim one party has against others that they do or refrain from some activity, and it is natural to think of these burdens as duties or obligations that are owed to the rights holder. If I have a right to free speech, then this seems to entail that others have a duty not to interfere with me in certain ways, and thus that the concept of duty must be used to explain what a right is. If so, then the concept of duty is more basic than that of a right, and so a rights-based theory is conceptually impossible.

Granted, it is common to explain the idea of a moral right of one party in terms of certain corresponding duties on the part of others. But as J. L. Mackie, a defender of rights-based moral theory, explains, instead of thinking of rights in terms of duties, "we could look at it the other way round: what is primary is *A*'s having this right in a sense indicated by the prescription 'Let *A* be able to do *X* if he chooses,' and the duty of others not to interfere follows from this."²²

So let us follow Mackie and suppose that there is no conceptual barrier to there being a rights-based theory of right and wrong action. How might it be developed? The idea would be to begin with a list of fundamental rights, including the right to life and various liberty rights, and then define or characterize the concepts of right and wrong action in terms of moral rights. Here, then, is how we might express the basic idea of right conduct for a rights-based theory:

- R** An action is right if and only if (and because) in performing it either (a) one does not violate the fundamental moral rights of others, or (b) in cases where it is not possible to respect all such rights because they are in conflict, one's action is among the best ways to protect the most important rights in the case at hand.

This principle—it is more of a *scheme*—all by itself is too abstract to be of any practical use. What needs to be added, of course, is a specification of *the* fundamental moral right or, in case there is more than one equally fundamental right, a specification of the most fundamental moral rights. Mackie proposes a single fundamental moral right: the right of persons to “choose how they shall live.”²³ But this right to choose is wide open, and to work our way from it to specific moral obligations, we will need to specify what sorts of claims people have in virtue of having this most basic right. Perhaps we can begin by recognizing the Jeffersonian moral rights to life, liberty, and the pursuit of happiness. And then, for each of these general rights, we might specify them further by recognizing a set of specific moral rights, including, for example, a right to free speech.

So, specifying a single fundamental moral right (or set of them) and working from what is fundamental to more specific moral rights is one task of a rights-based moral theory. However one works out the details of what the moral rights are, we need to keep in mind the obvious fact that in some contexts, moral rights will come into conflict. My right of free speech to yell “Fire!” in a crowded theater, for instance, may conflict with the rights to personal safety of other patrons in the theater. And this brings us to the issue of applying a rights-based theory to moral issues.

Applying a Rights-Based Moral Theory

Principle (or scheme) R purports to explain an action's being morally right (and by implication morally wrong) in terms of respecting fundamental moral rights. Clause (a) of R covers the easy case in which one's action simply does not come into contact with the moral rights of others. I get up in the morning and decide to eat Cheerios for breakfast. Unusual cases aside, this action has nothing to do with the moral rights of others—I'm morally free to eat the Cheerios or not. Clause (b), however, is where a rights-based approach to moral problems is most relevant: one can frame many of the disputed moral issues featured in this book as a conflict of rights: right to life versus right to choice; right to express oneself in speech and writing versus right to public safety; and so on.

So in applying R (supplemented with a theory of rights) to moral issues, the challenge is to find the best way of properly balancing competing rights claims in arriving at a moral verdict about what ought to be done. As I will explain more thoroughly later in connection with the ethics of prima facie duty, it is very doubtful that there is some fixed mechanical procedure that one can use in arriving at a correct or justified moral verdict in particular cases based on a consideration of competing rights. Rather, what one needs is what philosophers

call **moral judgment**—roughly, an acquired skill at discerning what matters the most (morally speaking) and coming to an all-things-considered moral verdict, where this skill cannot be entirely captured by a set of rules. The point to stress here is that, as with the other moral theories we are considering, applying a moral theory—its principles—to particular issues is not a mechanical process. But this does not take away from the value of such theories in guiding one's moral deliberations and subsequent choices.

Rights-Focused Approaches to Moral Issues

We have noted that talk of rights is very common in moral thought and discussion. However, we have also noted that in thinking about a moral issue in terms of competing rights claims, one need not accept a rights-based moral theory as just described. Consequentialists, Kantians, and natural law theorists can and do recognize rights—though on these theories rights are not what is most morally basic in the theory. So, because one may appeal to rights in discussing a moral issue without accepting a rights-based moral theory, we need to recognize what we may call **rights-focused** approaches to moral issues. To say that an author's approach to a moral issue is rights-focused is simply to say that the author appeals to rights as a basis for taking a stand on the issue at hand—the author may or may not also embrace a rights-based moral theory. In my article introductions, I have chosen to use the term “rights-focused” in summarizing the views of those authors who appeal primarily to rights in their article, unless the author makes it clear that he or she embraces a rights-based moral theory.

E. Virtue Ethics

Sometimes our moral thinking is dominated by thoughts about what sort of person one would be if one were to perform some action. The thought of living up to certain ideals or virtues of what a morally good person is like is crucial here. Being an unselfish person is an ideal that we may use in evaluating some course of action, and sometimes we may think, “Not helping her would be selfish on my part, so I'm going to help.” When our moral thinking takes this turn, we are evaluating actions in terms of virtue and vice. The ideas of virtue and vice have played a negligible role in the moral theories we have surveyed (at least as I have presented them).²⁴ However, inspired primarily by the ethical views of the ancient Greek philosophers Plato and Aristotle, **virtue ethics** makes the concepts of virtue and vice central in moral theory. Such theories, as I will understand them, take the concepts of virtue and vice to be more basic than the concepts of right and wrong, and thus propose to define or characterize the latter in terms of the former.²⁵

One might characterize right and wrong in terms of virtue and vice in different ways, but here (roughly) is how Rosalind Hursthouse, whose article on abortion is included in chapter 7, formulates a virtue ethical principle of right action:

- VE** An action is right if and only if (and because) it is what a virtuous agent (acting in character) would not avoid doing in the circumstances under consideration.

How are we to understand the concept of a virtuous agent featured in this principle? One straightforward way is to say that the virtuous agent is one who has the virtues. But what is a virtue? And which ones does the virtuous agent have?

A **virtue** is a trait of character or mind that typically involves dispositions to act, feel, and think in certain ways and that is central in the *positive* evaluation of persons. Honesty and loyalty are two commonly recognized virtues. The trait of honesty, for instance, involves at a minimum being disposed to tell the truth and avoid lying, as well as the disposition to have certain feelings about truth telling (positive ones) and about lying (negative ones). Honesty, as a virtue, is a trait that has positive value and contributes to what makes someone a good person. In contrast to a virtue, a **vice** is a trait of character or mind which typically involves dispositions to act, feel, and think in certain ways, and that is central in the *negative* evaluation of persons. So, for instance, opposed to the virtue of honesty is the vice of dishonesty which may be understood as having inappropriate dispositions of action and feeling regarding truth telling and lying. Furthermore, as a vice, dishonesty has negative value and contributes to what makes someone a morally bad person. So, in general, virtues and vices are character traits that are manifested in having certain dispositions to act and feel in certain ways and that bear on what makes a person morally good or bad. Here, then, is a short (and by no means complete) list of fairly commonly recognized moral virtues and their corresponding vices:²⁶

- Honesty/Dishonesty
- Courage/Cowardice
- Justice/Injustice
- Temperance/Intemperance
- Beneficence/Selfishness
- Humility/Arrogance
- Loyalty/Disloyalty
- Gratitude/Ingratitude

Applying Virtue Ethics

To apply VE to a particular case, then, we must first determine which character traits are the virtues that are possessed by the virtuous agent (we may begin with the previous list), and then determine how, on the basis of such traits, this agent would be disposed to act in the circumstances in question. An action that a virtuous agent, acting in character, would not fail to perform in some circumstance is morally required, an action she might or might not do at her discretion is morally optional, and one that she would avoid doing is morally wrong.

Of course, in applying virtue ethics to disputed moral issues, we encounter the fact that more than one virtue is relevant to the case at hand and that one of them—say, honesty—favors telling the truth, whereas one of the others—say, loyalty—favors telling a lie. In such cases of conflict among the virtues, we must examine the particular details of the case at hand and ask such questions as, “What is at stake here?” “How important is telling the truth in this case?” “How important is loyalty to an organization?” It is only by examining the details of such cases of conflict that we can come to an all-things-considered moral evaluation of some particular action based on considerations of virtue. This point is reflected in VE’s reference to a virtuous agent *acting in character*. Presumably, such an ideal agent has the sort of practical wisdom or judgment that is required in order for her to discern which virtue consideration, from among the competing virtue considerations in a particular case, has the most weight. As I have been noting all along in presenting the various moral theories—something that we explore a bit further in the next subsection—the application of moral theories to particular issues requires moral judgment.

F. Ethics of Prima Facie Duty

Whereas consequentialism, for instance, features a single moral principle of right conduct, what I am calling the **ethics of prima facie duty** features a plurality of basic moral principles of right conduct. The most famous version of this kind of view was developed by the twentieth-century British philosopher W. D. Ross (1877–1971). To understand the elements of Ross’s view, we need to do the following: (1) explain what he means by talk of “prima facie duty”; (2) present his basic principles of prima facie duty; and then (3) explain the role of moral judgment in applying them in practice.

The Concept of a Prima Facie Duty

To say that one has a **prima facie duty** to perform some action is to say that one has *some* moral reason to perform the action, but the reason in question might be *overridden* by some other moral reason that favors not performing the action. The best way to understand the concept is with an example. Suppose I have promised to pick you up on Saturday by 10:00 A.M. so that you can get to a very important job interview (your car is in the shop). Ross would say that because of my having made a promise to you, I have a prima facie duty (of fidelity—see later discussion) to do what I said I would do: pick you up by 10:00 A.M. on Saturday. But now suppose that as I am about to leave to pick you up, my child falls off the roof of my house and needs immediate medical attention. Ross would say that here I have a prima facie duty to take my child to the emergency ward of the hospital. So, I have a prima facie duty to start out for your place and a conflicting prima facie duty to attend to my child: as things have turned out, I am not able to fulfill both prima facie duties. Now the point of calling a duty “prima facie” is that the moral reasons provided by such facts as that I’ve made a promise or that my child needs my help can be outweighed by other moral reasons that favor doing some other action. Ross puts this point by saying that a prima facie duty can be overridden—beat out—by a competing prima facie duty. In the case I’ve described, because it is my child and because she needs immediate medical attention, my prima facie duty to help her overrides my prima facie duty to come pick you up. When one prima facie duty prevails in some conflict of duties situation, it becomes one’s *all-things-considered duty*—it is what you ought, all things considered, to do in that circumstance. So, for Ross, to say that one has a prima facie duty to perform action *A* on some occasion is to say that one has a moral reason to do *A*, and unless something comes up that is morally more important, one has an all-things-considered duty to do *A* on that occasion.

Ross’s theory of right conduct, which is our main concern, is based partly on his theory of intrinsic value to which we now turn.

Ross’s Theory of Intrinsic Value

Ross held that there are four basic intrinsic goods:

1. Virtue. The disposition to act from certain desires, including the desire to do what is morally right, is intrinsically good.
2. Pleasure. States of experiencing pleasure are intrinsically good.
3. Pleasure in proportion to virtue. The state of experiencing pleasure in proportion to one’s level of virtue is intrinsically good.
4. Knowledge. Having knowledge (at least of a nontrivial sort) is intrinsically good.

The items on this list are the basis for some of Ross's basic prima facie duties goods—call them “value-based” prima facie duties. What Ross calls duties of “special obligation” are not based on his theory of intrinsic value.

Ross's Prima Facie Duties

Here, then, is Ross's list of basic prima facie duties, organized into the two categories just mentioned:

Basic Value-Based Prima Facie Duties

1. Justice: prima facie, one ought to ensure that pleasure is distributed according to merit.
2. Beneficence: prima facie, one ought to help those in need and, in general, increase the virtue, pleasure, and knowledge of others.
3. Self-improvement: prima facie, one ought to improve oneself with respect to one's own virtue and knowledge.
4. Nonmaleficence: prima facie, one ought to refrain from harming others.

Basic Prima Facie Duties of Special Obligation

5. Fidelity: prima facie, one ought to keep one's promises (including the implicit promise to be truthful).
6. Reparation: prima facie, one ought to make amends to others for any past wrongs one has done them.
7. Gratitude: prima facie, one ought to show gratitude toward one's benefactors.

The first four basic prima facie duties, then, make reference to what has intrinsic value according to Ross's theory of value. Ross himself points out that the prima facie duties of justice, beneficence, and self-improvement “come under the general principle that we should produce as much good as possible.”²⁷ This part of Ross's theory fits the characterization of consequentialism.

The duties of special obligation do not make reference to what has intrinsic value; the duties of fidelity, reparation, and gratitude do not depend for their prima facie rightness on the values of the consequences of those actions. This part of Ross's theory is clearly duty-based or deontological. Overall, then, Ross's theory represents a hybrid: part consequentialist, part deontological.

Applying the Ethics of Prima Facie Duties

But how, on Ross's view, does one determine in some particular case that one prima facie duty overrides another, competing prima facie duty? Ross denies that there is any correct super-principle like the principle of utility or Kant's categorical imperative to which one might appeal to determine one's all-things-considered duty in cases of conflict. Nor is there any fixed ranking of the various prima facie duties such that the duty higher up on the list always beats out duties below it. Rather, according to Ross, in determining which prima facie duty is most “stringent” in some particular case and thus represents one's all-things-considered duty, one must examine the details of the case by using one's *judgment* about which of the competing duties is (in that situation) strongest. As mentioned earlier, moral judgment is a matter of discerning the morally important features of a situation and determining what ought or ought not to be done, where doing so cannot be fully captured in a set of rules. It is, therefore, a matter of skill that one may acquire through experience.

One final remark. One need not agree with Ross's own list of basic prima facie duties in order to accept the other tenets of Ross's view. For instance, Robert Audi has recently defended an ethic of prima facie duties that features ten basic prima facie duties.²⁸ Audi, unlike Ross, distinguishes duties not to lie from duties of fidelity, and he adds two additional duties to Ross's list. So were we to make the additions Audi proposes, we would have the following:

8. Veracity: prima facie, one ought not to lie.
9. Enhancement and preservation of freedom: prima facie, one ought to contribute to increasing or at least preserving the freedom of others with priority given to removing constraints over enhancing opportunities.
10. Respectfulness: prima facie, one ought, in the manner of our relations with other people, treat others respectfully.

The main point I wish to make here is that Ross's version of an ethic of prima facie duties is one version of this general sort of view. Audi's view attempts to build upon and improve Ross's view.

This completes our survey of some of the leading moral theories that figure importantly in many of this book's readings. As mentioned earlier, I recommend using these summaries of the six theories as an aid in understanding those writings in which an author appeals to one or other of them. The remaining section and the appendix address questions about moral theory that are likely to occur to the reader:

- What is the point of moral theory in thinking about disputed moral issues in light of the fact that there is a variety of competing moral theories?
- What about theories that appeal to the will of God or to the norms of society or culture in determining what is right or wrong?

3. COPING WITH MANY MORAL THEORIES

This chapter began with a brief overview of the central concepts and guiding aims of moral theory and then proceeded to survey six types of moral theory. In working through the various moral problems featured in this book, one will find that different moral theories often yield different and conflicting answers to questions about the morality of some action. The natural law theory, for instance, arguably condemns all homosexual behavior as morally wrong; a consequentialist approach does not. So the application of one theory to an issue may yield one moral verdict, while the application of another theory may yield a conflicting moral verdict. What, then, a student may ask, is the point of thinking about disputed moral issues from the perspective of moral theory? It all seems rather arbitrary.

This is a completely understandable question whose answer requires that one move from a focus on particular moral issues to questions about the nature and evaluation of moral theories. It is not possible to fully address such questions in a chapter whose aim is to provide students with a basic understanding of a range of moral theories. But because of its importance, the question does deserve to be addressed, even if briefly. In so doing, I will first offer some remarks about evaluating a moral theory, and then I will suggest some of the

the various moral theories for the illumination I think they provide in thinking about moral issues.

Evaluating a Moral Theory

Philosophers who develop a moral theory do not just state some moral principle or other and leave it at that; rather, they *argue* for whatever principles they are proposing. And we can critically evaluate their arguments. So the first point I wish to make is that there can be rational debate about a moral theory—not any old moral theory is as good as any other.

Furthermore, there are standards for evaluating a moral theory—standards that are not arbitrary but rather have to do with the guiding aims of a moral theory that we discussed in section 1 of this chapter. Corresponding to the theoretical aim of moral theory—the aim of explaining what makes something right or wrong, good or bad—is the principle of **explanatory power**:

A moral theory should feature principles that explain our more specific considered moral beliefs, thus helping us understand *why* actions, persons, and other objects of moral evaluation are right or wrong, good or bad. The better a theory's principles in providing such explanations, the better the theory.

This principle appeals to our “considered” moral beliefs, which may be defined as those moral beliefs that are *deeply held* and *very widely shared*. I hope that everyone reading this text believes that murder is wrong, that rape is wrong, and that child molestation is wrong. The list could be extended. Moreover, such moral beliefs are (for those who have them) very likely deeply held convictions. The principle of explanatory power tells us to evaluate a moral theory by determining whether its principles properly explain why such actions are morally wrong. Similar remarks apply to widely shared and deeply held beliefs about our obligations. So we can help confirm a moral theory by showing that it can properly explain the rightness or wrongness of actions about whose moral status we are virtually certain. Correlatively, we can criticize a moral theory by showing that it does not properly explain the rightness or wrongness of actions about whose moral status we are virtually certain. Applying this principle requires that we can tell what counts as a good explanation of the rightness or wrongness of actions. This is a topic of lively and ongoing philosophical inquiry whose study would take us far beyond the scope of this book. But in thinking about moral issues from the perspective of moral theory, the reader is invited to consider not only what a theory implies about some action or practice, but also what explanation it provides for whatever verdict it reaches about the action or practice under consideration. (I return briefly to this matter toward the end of this section.)

According to the practical aim of moral theory, we want moral principles that will help guide our moral deliberations and subsequent choices. Corresponding to this aim is the principle of **practical guidance**:

A moral theory should feature principles that are useful in guiding moral deliberation toward correct or justified moral verdicts about particular issues which we can then use to help guide choice. The better a theory's principles are in providing practical guidance, the better the theory.

Any moral theory that would yield inconsistent verdicts about some particular concrete action is obviously of no practical help on the issue at hand. Furthermore, a moral theory

whose principles are so vague that it fails to have clear implications for a range of moral issues is again of no help in guiding thought about those issues. Finally, a moral theory whose principles are extremely difficult to apply because, for example, applying them requires a great deal of factual information that is humanly impossible to acquire, is at odds with the principle of practical guidance. These are three measures to consider in evaluating how well a moral theory does in satisfying the principle of practical guidance and thus how well it does in satisfying the practical aim of moral theory.

These brief remarks are only meant to indicate how one can begin to evaluate a moral theory. Hopefully, what I have said is enough to make a start on answering the challenge that began this section. Let us now move on to the second point I wish to make in response to the challenge.

*Moral Theory and Moral Illumination*²⁹

I conclude with a plea for the importance of moral theory, even if there is no one theory that currently commands the allegiance of all philosophers who specialize in ethics. The plea is that moral theory can help focus and sharpen our moral thinking about particular issues, and it can thereby provide a kind of insight and illumination of moral issues that is otherwise easily missed. Let me explain.

No doubt readers of this chapter will have noticed that the various moral theories we have surveyed build on ideas that are very familiar. To see this, let us return to the case of euthanasia with which this chapter began. You may recall that in that case, Dr. Jack Kevorkian brought about the death of his patient Thomas Youk by a lethal injection. We described Kevorkian's action as an instance of voluntary active euthanasia. Now if one pays attention to on-line discussions and newspaper editorials that focus on this moral issue, and listens to the views of politicians and other social activists who discuss it, we find that some arguments appeal to the likely effects or consequences of allowing this practice. And of course, the idea that an action's rightness or wrongness is to be explained by reference to its likely consequences is the main idea of the various varieties of consequentialist moral theory. Similar remarks can be made about the other five types of moral theory presented in section 2. Some arguments over euthanasia focus on the intrinsic value of human life—one of the four basic human goods featured in natural law ethics. Related to questions about end-of-life moral decisions, some have argued that providing a terminal patient with painkilling drugs that will knowingly cause the patient to die of liver failure before succumbing to cancer is nevertheless permissible because death in this case is merely a foreseen side effect of the painkilling drug. Here we have a tacit appeal to the doctrine of double effect. Again, we find arguments that appeal to the special dignity and worth of human beings, as well as arguments that appeal to such alleged rights as the right to die or the right to die with dignity—arguments that tacitly appeal, respectively, to elements of Kantian moral theory and to rights-based moral theory (or at least rights-focused approaches to moral issues). Similar points can be made about virtue ethics and the ethics of *prima facie* duties.

So the first point I wish to make about studying moral issues from the perspective of moral theory is that one thereby gains greater insight and clarity into the kinds of arguments that one commonly reads and hears (and perhaps is disposed to give) over disputed moral issues. In fact, one may think of the various moral theories we have surveyed as attempts to develop such familiar ideas from moral thought and discourse in a rigorous philosophical manner. To really understand some moral issue for purposes of making up your own mind about it, you

first have to understand the issue, which in turn requires that you consider the various reasons that reflective people bring to bear in thinking and debating the issue at hand. Such reasons, as I have just indicated, are often developed systematically in a moral theory. So coming to understand moral theory helps provide a kind of moral illumination or insight into moral issues.

The further point is this. Different moral theories differ partly because of how they propose to *organize* our moral thinking about practical issues. For instance, utilitarianism has us organize our moral thinking about some issue in terms of its likely effects on well-being or happiness. Virtue ethics, by contrast, has us organize our moral thinking around considerations of virtue and vice, asking us, for example, to view a proposed course of action in terms of what it would express about our characters. Rights-based moral theories have us think about an issue in terms of competing moral claims that can be made by various involved parties. Similar remarks apply to Kantian moral theory, natural law theory, and the ethics of *prima facie* duty. But let us put aside for the moment the fact that the various moral theories in question purport to provide competing answers to questions about the underlying nature of right and wrong, good and bad. If we do, we might then view these theories as providing different ways of diagnosing and thinking about a moral problem, where in some cases the best approach is utilitarian, whereas in others the best approach is from a virtue ethics perspective, and still in others, some other moral theory best gets at what is morally most important to consider. In other words, it strikes me that some practical moral questions are best approached from, say, the perspective of act utilitarianism, others not. Here is an example that comes up in the chapter on war, terrorism, and torture and is discussed in the reading from Alan M. Dershowitz. He considers a “ticking bomb” scenario in which a captured terrorist very likely knows the whereabouts of a powerful explosive set to go off in a heavily populated city. Would it be morally permissible to torture this (uncooperative) individual in an attempt to extract information that might be used to locate and defuse the explosive? Given what is at stake in *this* particular case, I can well understand why one’s moral thinking would be guided by essentially act utilitarian reasoning. But in other cases, thinking in these terms seems morally askew. Thomas E. Hill Jr., in his article included in chapter 13, argues that in thinking about how we ought to relate to the environment, utilitarianism fails to properly diagnose what is wrong with certain ways of treating the environment. He also argues that thinking in terms of rights fails to get at what is really morally important about our dealings with the environment. His proposal is to think in terms of virtue—ideals of excellence—rather than in terms of utility or rights. As explained in section 1 of this chapter, a moral theory is partly in the business of providing practical guidance for moral thinking and decision making. My suggestion is that in some contexts it makes sense to think as an act utilitarian, in other contexts it makes most sense to think in terms of rights, and in still other contexts, thinking in terms of virtue and excellence seems most illuminating. The same can be said about the other moral theories we have surveyed. Thinking exclusively about all moral issues in terms of some one particular moral theory assumes a *one-size-fits-all* approach to moral thinking. I am suggesting that this probably isn’t the best way to use theory to illuminate practice.³⁰

Returning now to the challenge that began this section, I have tried to address it in two ways. First, moral theory is not arbitrary in the sense that you can just pick and choose your favorite or make up your own: there are standards for evaluating moral theories that have to do with the theoretical and practical aims of moral theory. Second, the variety of moral the-

ories on offer can positively aid in one’s moral thinking about controversial moral issues in two ways. First, it can do so by providing rigorous articulations of common ideas about morality. And second, it can do so if one views these theories as diagnostic tools for getting to the heart of moral problems. Some tools are better for some jobs than other tools. My suggestion is that a particular moral theory may be a better tool than others when it comes to thinking through some particular issue, though a different theory may be better at helping one think through other issues.

APPENDIX: ETHICS BY AUTHORITY?

DIVINE COMMAND THEORY AND ETHICAL RELATIVISM

The idea that morality depends on some authority—whether the will of God or the norms of one’s culture—is well known, even if not generally well understood. The readings in this collection by and large reflect the impact of the six moral theories presented in section 2 on philosophical thinking about disputed moral issues. This does not mean that other moral theories or approaches to moral issues are not worthy of philosophical attention. My presentation has been selective. However, because what I am calling ethics by authority in one form or another will likely occur to readers, I think it is important to explain why many who think about moral issues have grave reservations about both divine command theory and ethical relativism.

According to the **divine command theory**, what is right or wrong depends on God’s commands in the sense that what *makes* an action right or wrong are mere facts about God’s commands, nothing more. On this view, an action is wrong whenever (and because) God commands that we not do the action. An action is morally obligatory whenever (and because) God commands that we do it. Otherwise an action is morally optional. So the fundamental moral principle of this sort of theory can be expressed as follows:

DCT An action is right if and only if (and because) God does not command that we not do that action.

For many people, being told that God does or does not command some action is crucial in their thinking about moral issues. But surely if God commands that we perform some action, there must be some reason why God issues this command—some reason that explains *why* the action is something we ought or ought not to do. But then, as philosophers, we can ask what it is about the action in question that makes it wrong and is a basis for God’s command. And once we put the question to ourselves in this way, we are simply raising the general moral questions about the right and the good that we began with. So, appealing to God’s commands (at least for believers) may help the believer decide what to do, but the fact that God commands this or that action does not answer the deep question about the underlying nature of right and wrong that a moral theory attempts to answer—it does not plausibly address the main *theoretical aim* of moral theory explained earlier.

Of course, someone sympathetic to DCT may claim that it is just God's commands that make an action right or wrong. But this won't be acceptable to a theist who thinks that God's actions are rational. After all, if one says that there is no reason behind God's commands, then one is saying that God has no good reason for commanding that we keep our promises and not commit murder, that God's commands are completely arbitrary. But this can't be right. So, a theist must say that there are facts about an action that make it wrong and that since God knows all facts, and since God is all-good, God commands that we do what is (independently of his commands) right and not do what is wrong.

As for **ethical relativism**, there is a good deal of confusion generated by the vague (and unfortunately popular) talk of morality being relative. Surely anyone can agree that whether a particular action—say, addressing a professor by her first name—is morally right may be importantly affected by what a society considers to be insulting. In the United States at present, the social norms that help specify what constitutes an insult do not seem to consider a student addressing a university professor by her first name as an insult. If that's right, then a student in a U.S. university would not be insulting a professor in addressing her by her first name (unless, perhaps, the professor had expressed a desire not to be so addressed). But in other countries (at present), the social norms governing student-professor relationships are such that the sort of address in question does constitute an insult. If we agree that insults are morally wrong (see the chapters dealing with hate speech and with sexism, racism, and reparation for more about the morality of insults), then we can easily see that the action of addressing a professor by her first name is morally wrong in some social circumstances (when in certain countries) and not in others (when in the United States). But this kind of context-sensitivity of morality according to which one's circumstances, including the social norms of one's culture, may have a bearing on what is right or wrong to do in that culture is something that all of the moral theories we have considered do accept.

So if ethical relativism is to represent a moral theory that competes with the ones we have surveyed, what must it say? It must say something like this: (1) there are no correct moral norms or principles that are valid for all cultures at all times; rather (2) there are only the moral norms that some group or culture happens to accept, and these norms—*no matter what those norms say*—are what determine what is right or wrong for members of that group or culture. We can encapsulate these ideas in the following principle:

ER An action (performed by members of a group G) is right if and only if the moral norms that are accepted by G permit the performance of the action.

Thus, if some culture accepts the moral norm that the enslavement of other human beings is morally right, then (according to the relativist) enslavement really *is* right—for those people.

Now relativism has its popular allure. Some people seem to take ethical relativism as an enlightened view about the true nature of morality. In order to disabuse the reader of the kind of simple ethical relativism just described, consider abortion. Suppose we find out that a majority of current U.S. citizens accept a set of moral norms that find nothing wrong with abortion. If we suppose that the moral norms of some culture are those norms that are held by a majority of its members, then according to ethical relativism, we would have to conclude that abortion (for members of U.S. culture) is morally right. Even if you think that abortion is morally right, do you (the reader) really think that the actual moral rightness and wrongness of an action depends on majority opinion? If so, then you might think we can settle moral issues by a vote. But this would mean that no matter what the majority of some group

accepts as part of that group's moral norms—genocide, slavery, infanticide, lying, cheating, whatever—those actions would be right for members of that group. Granted, the members of some group may honestly *think* that genocide is morally right, but thinking something is right does not *make* it right. Right?

So, we can agree to the following thesis of **context-sensitivity**:

CS The rightness or wrongness of an action may depend in part on facts about the agent and her circumstances, where her circumstances may include facts about the norms for what counts as constituting insults, a person's privacy, proper respect for others, and so forth.

The example of insulting behavior illustrates how CS can be true. But as we have already noted, CS is compatible with all of the nonrelativist moral theories that we have surveyed in section 2. According to each of those theories, there are basic moral principles or norms whose correctness is objective and not dependent on whether they happen to be accepted by some culture. Of course, in applying one of these principles to some particular case, we must consider various details of the case including facts about particular agents and their circumstances as well as facts about the society in which one happens to live.

Whether some version of ethical relativism can be defended is controversial and cannot be settled here. If so, it would have to improve upon ER. My main point was to note the important difference between CS and ER (they are often confused) and explain why CS is uncontroversial and why ER is problematic.

According to both divine command theory and ethical relativism, morality depends on the dictates of some authority—God or culture. I have tried to indicate very briefly why many moral philosophers are not satisfied with either of these theories.³¹ In any case, in our readings, although a few authors appeal to theological premises in attempting to support a position on a disputed moral issue, none of these authors rest their case solely on a brute appeal to what they take to be God's commands. And none of the authors appeal to ethical relativism.

NOTES

1. A few paragraphs of material in this essay are taken from my "Ethics" in *Reflections on Philosophy*, 2nd ed., ed. L. McHenry and T. Yagisowa, (New York: Longman's Publishers, 2003), 103–25.

2. John Rawls, *A Theory of Justice* (Cambridge, MA: Harvard University Press, 1971), 24.

3. Given this understanding of the notions of intrinsic and extrinsic value, it is possible for something to have value of both sorts. Suppose, for example, that both happiness and knowledge have intrinsic positive value. Since knowledge can be of use in promoting happiness, knowledge can also have extrinsic value.

4. The "if and only if (and because) . . ." is meant to make clear that what follows the "and because" is meant to be a moral criterion that explains *why* the item being evaluated has whatever moral property (e.g., rightness) is mentioned in the first part of the principle.

5. In section 2D the idea of a rights-based moral theory will be introduced.

6. See Jeremy Bentham, *An Introduction to the Principles of Morals and Legislation* (New York: Hafner Press, 1948, originally published in 1789), and J. S. Mill, *Utilitarianism* (Indianapolis, IN: Hackett Publishing, 1979, originally published in 1861).

7. Another important distinction within consequentialism is between versions that appeal to the *actual* consequences (and associated value) that would occur were some action to be performed and versions that appeal to the *likely* consequences (and associated value) of actions—those consequences and their associated value that can be reasonably expected to follow from an action were it to be per-

formed. I have chosen to formulate consequentialist principles in terms of likely consequences, since in applying a consequentialist theory to practice, we have to rely on our *estimates* of the consequences and associated value of actions.

8. I have explained utilitarianism in terms of *human* happiness or welfare, but a utilitarian may expand the scope of moral concern to all creatures for whom it makes sense to talk about their happiness or welfare.

9. For a recent defense of perfectionist consequentialism, see Tom Hurka, *Perfectionism* (Oxford: Oxford University Press, 1993).

10. What I am describing is a pure perfectionist account of value. It is, of course, possible to accept a hybrid view of intrinsic value according to which both happiness and perfectionist goods such as knowledge and achievement have intrinsic value.

11. For a recent defense of rule consequentialism, see Brad Hooker, *Ideal Code, Real World* (Oxford: Oxford University Press, 2000).

12. This is not to say that the best form of the natural law theory embraces the idea that “unnatural” actions are wrong and “natural” actions are right. The version I am about to present does not feature such ideas.

13. For a recent defense of natural law theory, see John Finnis, *Natural Law and Natural Rights* (Oxford: Oxford University Press, 1980).

14. Here is an appropriate place to clarify what I am calling theories of intrinsic value that figure importantly in those moral theories which feature value-based theories of right conduct. Ideally, a complete theory of intrinsic value would accomplish two related tasks: (1) provide a complete list of those types of things that have intrinsic value and also (2) specify those underlying features of intrinsically good and bad things *in virtue of which* they have whatever intrinsic value they do have. However, since our main concern in this chapter is with that part of a moral theory having to do with right and wrong action, we need only consider how a theory of intrinsic value responds to the first task—a specification of the types of things that have intrinsic value. So, for example, in what immediately follows, I will simply list the most basic types of items that have positive intrinsic value according to the natural law theory.

15. Kant’s major writings in ethics include *Groundwork of the Metaphysics of Morals* (1785), *Critique of Practical Reason* (1790), and *The Metaphysics of Morals* (1797). All of these writings are included in Mary E. Gregor, trans., *Kant’s Practical Philosophy* (Cambridge: Cambridge University Press, 1997). Page references to Kant’s writings are to this volume. For a recent defense of Kantian ethics, see Onora O’Neill, *Towards Justice and Virtue* (Cambridge: Cambridge University Press, 1996).

16. In her article on a Kantian approach to world hunger included in chapter 11, Onora O’Neill develops these ideas in more detail.

17. See *Groundwork of the Metaphysics of Morals*, section II, 84–85.

18. I have left out the “and because” since arguably this formulation does not purport to express a moral criterion of right action—what *makes* an action right for Kant is expressed by the Humanity formulation; a *test* of an action’s rightness is provided by the Universal Law formulation. For more on this, see my *Moral Theory: An Introduction*, chap. 5.

19. In the *Groundwork*, section II, in the footnote on p. 80, Kant raises objections to the golden rule.

20. *Practical Philosophy*, 74.

21. Strictly speaking, rights-based moral theory would construe rights as more basic than right action (including duty) and value. But here we need only be concerned with moral theories that take rights to be more basic than right action.

22. J. L. Mackie, “Can There Be a Rights-Based Moral Theory?” *Midwest Studies in Philosophy* 3 (1978): 351.

23. Mackie, “Rights-Based Moral Theory,” 355.

24. This does not mean that such theories have little to say about such matters. For instance, Kant elaborates a theory of virtue in the “Doctrine of Virtue” which makes up part 2 of his 1797 *Metaphysics of Morals*.

25. For a recent defense of virtue ethics, see Rosalind Hursthouse, *On Virtue Ethics* (Oxford: Oxford University Press, 1999).

26. Defenders of virtue ethics often attempt to explain the basis of the virtues and vices (why some trait is a virtue or a vice) by appealing to the idea of human flourishing. The idea is that a trait of character or mind is a virtue because it contributes to or partly constitutes the flourishing of its possessor. For more on this point, see Hursthouse’s “Virtue Theory and Abortion” included in chapter 7.

27. W. D. Ross, *The Right and the Good* (Oxford: Oxford University Press, 1930), 27.

28. See Audi’s *The Good in the Right* (Princeton, N.J.: Princeton University Press, 2004) for a recent defense of an ethic of prima facie duty that attempts to integrate this sort of view into a basically Kantian framework.

29. Special thanks to Jason Brennan and to Dave Schmidtz for very helpful conversations about moral theory and illumination.

30. These remarks suggest the possibility of combining certain elements from the various theories into one big super-theory featuring a plurality of principles, some having to do with duties, others with virtuous actions, others with rights, others with utility, and perhaps all of them unified by the Kantian idea of respect for persons, animals, and the environment. Doing so would still leave open the question of whether some one element of the theory—duties, virtues, rights, etc.—is most basic in the theory. One possibility (and the one that strikes me as initially plausible) is a theory according to which these notions are “interpenetrating”—a full understanding of any one of them requires appeal to the others.

31. For a more detailed presentation and critique of divine command theory and ethical relativism, see chapters 2 and 3 of my *Moral Theory: An Introduction* (Lanham, MD: Rowman and Littlefield, 2002). This book also has chapters on utilitarianism, natural law theory, Kantian ethics, the ethics of prima facie duty (what I there call ethical pluralism), and virtue ethics.