Discussion Topic Guidelines

1. Assignment: Twice, in teams of five, your assignment is to become informed about a moral issue in technology, which we will informally discuss on Mondays during unit two. Your team will be evaluated based on two components—a written and an oral component, as follows:

(a) The written component: Your team will write a roughly two page informal report on your assigned topics (500-700 words), to be turned in at the beginning of class on the day of our discussion. Late papers will not be accepted. Roughly, your instructions are as follows: You will briefly introduce the issue, and then provide some moral reasons both for and against.***** More specific instructions for your particular prompt can be found on the pages below.

> ***** Note: For the purposes of our course, when brainstorming reasons for and against each position, the focus of your efforts should be on the *moral* reasons for and against; though you *may* also appeal to practical and/or legal reasons—especially in instances where you are able to make a case that these other sorts of reasons might be morally relevant.)

(b) The discussion component: On the assigned date, your team will lead a class discussion on that topic. I will come prepared with a set of questions and comments, which I may interject with periodically. (This will become more likely if I see that the discussion is flagging.)

2. How to Begin: Begin by reading your particular prompt (on the pages below). I then encourage you to do some preliminary brainstorming and research on your own. As you begin to familiarize yourself with the issue and read articles about it, ask yourself: What is *my* moral stance on this issue? *Why* do I believe this? Why do others disagree? What reasons do they have for their stance?

You should then arrange to meet up with your teammates to discuss and share your thoughts from your preliminary research with one another. From there, you can then do some further brainstorming as a group, and decide how to divide up the work, and so on.

3. Grading Rubric: I will be looking for several things when I assign grades:

Written

(1) Following Instructions: Did you complete the assignment according to the instructions stated here?

(2) Clarity: Do you explain yourself in a way that is *clear*, *concise*, and *well-organized*? You should think of this as a somewhat informal assignment. However, your writing should still be clear enough and careful enough that someone who has never taken philosophy could read it and understand it (and perhaps even learn from, and be persuaded by it!).

(3) Careful, Contemplative Reasoning: It should be evident that you have actually thought carefully about the assigned issues, and that you have put some time and consideration into your written response. In short, this is not meant to be the sort of assignment that can be completed successfully 10 minutes before class.

Oral

(1) Preparedness: It should be apparent that each of your team members has come to class prepared to discuss the issue—in both your prepared presentation as well as your responses to any questions and comments.

(2) Class Engagement & Facilitation of Discussion: Ask yourself, What sort of classroom experience do *you personally* value the most? Do you prefer for class to be fun? to raise interesting and thought-provoking questions? to provide helpful and insightful potential answers to those questions? to be engaging? and encouraging of discussion? 'Yes' to all of the above? Then strive for that. Do not plan to deliver an uninterrupted lecture; rather, plan to come to class with the intention of simply introducing an issue, and then helpfully guiding a fun and productive discussion about it.

(3) Civility: It should go without saying that our discussion will remain civil and respectful. This means no insulting of classmates, or shouting at or over them, and it also means giving others the opportunity to share their own views. We will also strive whenever possible to keep our comments constructive and productive, with the goal of moral progress and learning as we work through these difficult issues together, in a group effort.

4. Specific Topics: Specific prompts for the six discussion topics can be found on the pages below.

Day One (Mon, 10/5): Facial Recognition

The Issue: Facial recognition technology is here. A brief timeline of some notable events:

- August 2017: Police use the tech to identify and catch an MS-13 gang member (here) and later (June 2018) a mass shooter (here).
- September 2017: Apple launches 'Face ID'. (Users can now use their faces to unlock their phones, log in to apps, and authenticate purchases.)
- February 2018: Police officers in China begin wearing glasses with facial recognition abilities (here).
- May 2018: Taylor Swift uses facial recognition at her concert in order to cross-check audience members against a database of her known stalkers (here).
- Fall 2018: Starting in Seattle (here), and increasingly elsewhere (here), schools across the country begin using the technology to restrict entry to school grounds.
- November 2018: The Secret Service begins testing the tech around the White House (here); ICE explores the use of this technology to identify undocumented immigrants (here).
- December 2018: Delta launches the nation's first facial recognition airport terminal, where passengers show their faces rather than their ID's (see here or here). Since then, the technology has rapidly been implemented in a number of other terminals, especially for international flights (more info here). The tech was already being used in airports by CBP (Customs and Border Protection), whose 'biometric verification' program caught its first passport imposter in August, 2018 (here).
- May 2019: San Francisco becomes the first of a growing number of U.S. cities to ban the tech (here).
- December 2019: China implements mandatory face scans for all new mobile devices (here).
- End 2020: China plans to have fully launched its nationwide mass surveillance system, including over 600 million surveillance cameras with facial recognition (here, here) and a social credit system (here).

It's well past time to start considering the moral implications of this technology. For instance, do we as human beings have a fundamental right to privacy; and, if so, does the use of this technology violate that right? Does using this technology make us safer, or better off in some other way; and, if so, do these benefits outweigh any potential costs?

Question: Should we pass legislation to prohibit public use of facial recognition technology?

For some excellent discussion of this issue, listen here. Some videos here, here, and here.

For the Writing Assignment: Briefly introduce the controversy surrounding facial recognition technology, and then discuss at least **two reasons in favor** of prohibiting public use of it, and **two reasons against** such a prohibition. Then, decide which side of this issue your team ultimately agrees with, and explain why.

Day Two (Mon, 10/12): Social Media Content (or, Facebook, Twitter, & Censorship)

The Issue: As social media platforms become increasingly more powerful (now our primary source of news, communication, etc.), questions regarding their responsibility for, and obligations with respect to, the content that appears on their platforms become ever more pressing.

A brief timeline of some notable recent events

- *2016:* Social platforms such as Facebook are widely used in order to spread disinformation in an effort to affect the results of the 2016 presidential election (e.g., here and here).
- *September 24, 2019:* Facebook announces that it will not fact-check or censor political ads (here).
- *October 17, 2019:* Facebook CEO Mark Zuckerberg addresses the controversy by speaking at Georgetown University in favor of free speech (here).
- *October 23, 2019:* Alexandria Ocasio-Cortez presses Zuckerberg on the issue before Congress (here).
- *January 9, 2020:* Facebook reaffirms that it will not fact-check or censor political ads (here).
- *May 2020:* Facebook and YouTube repeatedly scrub all instances of 'Plandemic' documentary videos, for spreading misinformation about COVID-19 (here).
- *May 26, 2020:* Twitter takes action against President Trump's tweets for the first time, labeling a two-part tweet about the fraud associated with voting by mail (here and here) as potentially misleading.
- *May 28, 2020:* President Trump issues an executive order, forbidding Facebook, Twitter, Instagram, YouTube, etc., from censoring content in the form of deleting it or flagging it as inappropriate or misleading (here).
- *May 29, 2020:* Twitter ignores executive order, flagging the second of another two-part tweet from President Trump regarding the George Floyd protests (here) for glorifying violence.
- *June 18, 2020:* Twitter flags Trump's "racist baby" video re-tweet (here); Twitter, Facebook then remove video. That same day, Facebook removes Trump campaign ad for using Nazi symbol (here). (Story about both here.)
- *June 23, 2020:* Twitter flags another Trump tweet directed at DC protestors (here) for abusive behavior.
- *July 27, 2020:* Facebook, Twitter, YouTube scrub all instances of video of press conference held by Dr. Stella Immanuel and others (Facebook and Twitter also removing President Trump's posts and retweets of the video), for spreading misinformation about COVID-19 (here , here).
- *July 28, 2020:* Twitter removes tweet, which Trump had re-tweeted, for falsely claiming that there is a cure for COVID. (here)
- *August 5, 2020:* Facebook and Twitter remove videos of President Trump claiming that children are "almost immune from this disease" (i.e., COVID-19), and temporarily suspend his campaign's account. (story here)
- *August 23, 2020:* Twitter flags a Trump tweet (here) for spreading misinformation about voting by mail.
- *August 31, 2020:* Twitter removes Trump re-tweet of claim that only 6% of reported coronavirus deaths are real. (here) That same day, they also flag the Trump campaign's tweet of a video of Joe Biden for being "manipulated media". (here)

The nation is split over what the role of social media platforms should be with respect to its content. Should these platforms *police* their content? Or, do we have a fundamental right to freedom of speech which such content-policing would violate?

Question: Should social media platforms be moderating or policing user-generated content? If so, what kind of content should be censored? For instance, is it morally permissible (or even obligatory) to delete or flag a post for merely containing false information? Why or why not? Should the president's posts be granted an exception to such a policy? Why or why not?

See this article (here) and brief video (here) for some further discussion of the issue.

Related issues: Video synopsis of the earlier controversy over political ads here; NPR story on ads here; and a humorous study of the spread of conspiracy theories on the internet here.

For the Writing Assignment: Briefly introduce the issue, and then discuss at least **two reasons in favor** of social media platforms censoring or flagging misleading content, and **two reasons against** their doing so. Then, decide which side of this issue your team ultimately agrees with, and explain why.

During the course of your discussion, be sure to also explore the question of whether publicly elected officials (e.g., Trump) should be held to the same standards as others, as well as the broader question about what *sorts* of content should be censored, specifically (if any).

Day Three (Mon, 10/19): Predictive Algorithms (or, Robot Criminal Justice)

The Issue: Today, AI's and predictive algorithms are helping to do the work of lawyers (here), judges (here), and police officers (here). For example, police departments across the country are making use of predictive algorithms (like PredPol or Hunchlab) to analyze large data sets and then predict where the next crimes will be, so that they can better patrol those areas; and judges make use of similar software (e.g., Compas/Equivant) to help them to make decisions about whether (and by how much) to let someone out of jail on bail before trial. These days, you can even hire an AI lawyer (e.g., Ross or LawGeex).

There is no doubt: Artificial intelligence is already a part of our criminal justice system. But, the question is, *Should it be?* Those using the software claim that it is helping to reduce crime. Others object to the software on grounds of racial bias and lack of transparency.

Question: Is the use of A.I.s and predictive algorithms in the criminal justice system morally permissible? For instance, is the use of so-called "predictive policing" software permissible?

Great introductory video here; excellent article in *The Atlantic* here (audio version here). Some further discussion here. See also here, here, and here. For a comedic take, see here.

For the Writing Assignment: Briefly introduce the issue, and then discuss at least **two reasons in favor** of the use of artificial intelligence in the criminal justice system, and **two reasons against**. Then, decide which side of this issue your team ultimately agrees with, and explain why.

Day Four (Mon, 10/26): Digital Contact Tracing

The Issue: Here we are in the midst of a global pandemic that will have claimed roughly one million lives by the end of 2020. Pretty terrible, right? Not to mention all of the other bad stuff like quarantine, social distancing, masks, and the end of social gatherings. But, life could return to normal more quickly with large-scale testing and contact-tracing efforts.

For instance, if you tested positive for COVID-19, it would be great if we then had some way of (a) knowing exactly with whom you've been in contact with recently, and (b) contacting them right away to let them know of their potential exposure. Enter "digital contact tracing", the proposal to use smartphones to track your whereabouts and contacts for the purposes of contact tracing and notification.

Interestingly, in August, **Virginia became the U.S. state to deploy the Apple/Google contact tracing app** (see here). It works by installing a setting on your smartphone that keeps track of your contacts by pinging the other phones around you via Bluetooth (see here). This sort of tech had already been implemented in Singapore in their TraceTogether program—supplemented by a bluetooth device given to those who do not have smartphones (here)—though only about a third of the population opted into Singapore's program, sparking questions of whether it should be mandatory (here). What do you think?

Some people worry that such a program would be a violation of our fundamental right to privacy, or that it might grant the government access to too much information and control of our lives. On the other hand, the program could potentially save tens or even hundreds of thousands of lives.

Question: Is it morally permissible for the government to implement a digital contact tracing program during a pandemic?

Some introductory videos on this topic here, here, and here. John Danaher's (yes, the sexbot guy) interview with Carissa Véliz on this issue here. Some further discussion here and here.

For the Writing Assignment: Briefly introduce the issue of digital contact tracing, and then discuss at least **two reasons in favor** of its use, and **two reasons against**. Then, decide which side of this issue your team ultimately agrees with, and explain why.

Along the way, be sure to consider some of the various other moral questions in this vicinity that might also arise. For instance, If not enough people opt into the program (as happened in Singapore), should the use of the application be made *mandatory*? Would it be permissible to extend the use of such a program beyond contact tracing? For instance, would it be permissible to use it to monitor (or even enforce) social distancing, or quarantining? Imagine for example that W&M implemented a program where every student was required to install location and contact tracking on their phones. Then, anyone detected at a gathering of more than ten people, or anyone detected violating quarantine after testing positive for the virus, etc., would be subject to academic sanctions. Would a program like that be morally permissible? Why or why not?

Day Five (Mon, 11/2): Resource Allocation (or, Who Gets the Ventilator?)

The Issue: Here we are in the midst of a global pandemic that will have claimed roughly one million lives by the end of 2020. Pretty terrible, right? Worse still, resources continue to be scarce. For instance, in some cases there are not enough ICU beds, ventilators, therapeutic drugs (e.g., Remdesivir), or simply health care *workers themselves* to go around to all of those who need them. In such cases of shortage, the question arises: Who gets the treatment, and who doesn't?

Question: What is the best metric or method for determining who gets medical treatment during a pandemic in cases of a shortage of resources? i.e., who gets the ventilator, and why? What is the fairest or most just distribution of life-saving resources?

See this *New York Times* article for an introduction to the issue (here). Polls showed that Americans are split when asked to choose between giving ventilators to 'those most in need at the moment' and 'those most likely to recover' (here). But, that poll ignored the host of other popular proposals, such as 'those who are youngest' or even 'whoever wins a coin toss'. For a great, brief overview of some of the major answers to this question, see here (which cites the authors of this longer article here). See also John Danaher's (yes, the sexbot guy) interview with Lars Sandman on this issue here.

For the Writing Assignment: Briefly introduce the issue of resource allocation, and then discuss and explore **at least two different proposed metrics** for allocating resources, citing at least **one reason for** and **one reason against** each of them. Then, decide what your team's final stance on this issue is, and explain why.

Along the way, be sure to consider some of the various other moral questions in this vicinity that might also arise. For instance, what would your team say about a future *vaccine* to the virus—which is sure to be in short supply when it first arrives? Who gets the vaccine first? Why?

Or, consider: When allocating, does it make a moral difference if the technology is already *in use* (e.g., *removing* a ventilator from person A who is already using it to give to person B, versus simply choosing to hook up a presently unused ventilator to person B rather than A)?

There are also considerations regarding the *globally* just distribution of resources. (Example: Was it permissible for the U.S. to buy up the entire global supply of Remdesivir? See here.) Is it permissible to significantly raise the price of life-saving devices, medications, or equipment when they are scarce? (Example: Is it permissible to sell N-95 masks for thousands of dollars? See here.) Does your team's conclusion about the fairest or most just distribution of resources during a pandemic also apply to *non-pandemic* decisions; or are our decisions about resource allocation morally different than at other times? And so on.

Day Six (Mon, 11/9): Outsourcing Risk (or, Using Instacart During a Pandemic)

The Issue: Here we are in the midst of a global pandemic that will have claimed roughly one million lives by the end of 2020. Pretty terrible, right? At least we don't have to go out for food. Why take the risk yourself, when you can outsource that risk to someone else? This year, use of the already popular host of food delivery apps (e.g., Instacart, DoorDash, Uber Eats, Grubhub, Postmates, etc.) has gone through the roof. But, is the use of these apps during a pandemic morally permissible?

Essentially, by using such an app rather than going out yourself, aren't you just getting someone else to take on your own risk for you? Is there anything morally objectionable about doing this? For starters, does it matter, morally, that the majority of those *using* the apps are (on average) among the best off, economically, while those *working for* the apps are (again, on average) among the worst off, economically? Or, is hiring someone to do something always permissible so long as they *agree* to do it?

Question: Is it morally permissible to use a food delivery app (e.g., Instacart) during a pandemic?

For further discussion, see here. Or watch Harvard philosophy professor Michael Sandel lead an excellent discussion with his students about it here (section from 28:36 - 44:41)

For the Writing Assignment: Briefly introduce the issue, and then discuss at least **two reasons in favor** of the permissibility of using food delivery apps during a pandemic, and **two reasons against**. Then, decide which side of this issue your team ultimately agrees with, and explain why.

Along the way, be sure to consider some of the various other moral questions in this vicinity that might also arise. For instance, what is the nature of consent, and how is it morally relevant to this issue, if at all? What is the nature of exploitation; and is it ever permissible? Would your moral conclusions be altered if there were no money exchanged? (for instance, if Instacart workers were *volunteer-only*?)

Also, might your conclusions about food delivery apps extend to *other* instances of outsourcing risk? For instance, when we discover a vaccine to the virus, it will need to be tested on human beings before widespread dissemination. Is it permissible to pay someone to participate in a clinical trial of a medication that may have as-yet-undiscovered negative health side-effects? Does the outsourcing of risk in this instance differ morally from that of using, say, Instacart? Why or why not?

Furthermore, be aware that, independent of the pandemic there was already moral controversy surrounding these apps on other grounds. For instance, apps like Uber Eats, Grubhub, and DoorDash have been criticized on the grounds that they harm local businesses by taking a huge chunk of their profits (see here and here); while all of them have been criticized for the poor treatment of their workers—e.g., low pay and lack of benefits (here).